

Eduardo Adame

Exact Computation for the Normalised Power Prior

Rio de Janeiro

2026

Eduardo Adame

Exact Computation for the Normalised Power Prior

Master's Thesis presented to the Graduate Program in Applied Mathematics and Data Science of the School of Applied Mathematics as a requirement to obtain the degree of Master in Applied Mathematics and Data Science.

School of Applied Mathematics

Fundação Getulio Vargas

Graduate Program in Applied Mathematics and Data Science

Supervisor: Luiz Max Carvalho, PhD

Co-supervisor: Flávio B. Gonçalves, PhD

Rio de Janeiro

2026

STATEMENT OF AUTHORSHIP

I hereby declare that the thesis submitted is my own work. All direct or indirect sources used are acknowledged as references. I further declare that I have not submitted this thesis at any other institution in order to obtain a degree.

To anyone who has ever believed in Education as the main way to overcome adversity.

To the memory of Marcos Vinícius Gomes, the funniest friend I could ever have had.

Acknowledgements

Out of respect for those I wish to thank, these acknowledgements are written in Portuguese, my native language.

Não seria possível começar esses agradecimentos sem mencionar meus pais, Patricia e Jeferson, e meu irmão Gabriel, que fizeram o seu melhor para me apoiar ao longo dos anos. Sei que não foi fácil acompanhar meus passos de longe, mas sou muito grato pela confiança que depositaram em mim. Agradeço também ao Leandro e à Kênnia por acompanhá-los durante o dia a dia; vocês foram essenciais. Tudo por vocês ainda é pouco.

Ao Caio Lins e ao Tulio Koneçny, meus queridos amigos e companheiros de apartamento: carregarei com muito carinho todas as memórias e registros que fizemos juntos. Obrigado por todas as gargalhadas e por tudo que me ensinaram. Agradeço também por embarcarem em todas as minhas maluquices, mesmo eu não sendo a pessoa mais fácil do mundo. Seremos sempre a Bald Comics.

Aos meus amigos Ezequiel, Carol, Pablo, Isaque, Juliano, Guilherme e Jairon, obrigado por todas as interações especiais que tivemos juntos.

Aos professores Yuri Saporito, Rodrigo Targino, Diego Mesquita, Hugo de la Cruz e Marcelo Sant'Anna, deixo meu agradecimento por todos os conselhos, pelo conhecimento compartilhado e, sobretudo, pelo suporte nas decisões mais importantes da minha trajetória.

Aos meus co-autores em outro trabalho, Guilherme T. Goedert e Daniel Csillag, obrigado pelo suporte e pela colaboração que tornaram possível construir um trabalho de qualidade num espaço com tanta visibilidade como o NeurIPS.

À Cássia Pessanha, agradeço o carinho e o cuidado que teve comigo, sobretudo na ausência de familiares por perto. Aos professores César Camacho e Maria Izabel Camacho, agradeço a oportunidade de ingressar na EMAP através do FGV CDMC, e ao Professor César, também, por toda a orientação e apoio à frente da direção da escola. Aos funcionários Ivone, Mônica, Marcio, Jorge, Ingrid, Roseni, Thais, João, Isabel e Cristiane, obrigado pela dedicação ao funcionamento da nossa instituição. Aos meus alunos de monitoria e aos professores responsáveis pelas disciplinas, obrigado pela confiança e pelo aprendizado mútuo. Ensinar é, sem dúvida, uma das melhores formas de aprender.

Por fim, ao meu co-orientador Flávio Gonçalves, que me instigou a reflexões fundamentais na minha jornada, e ao meu orientador Luiz Max Carvalho, que nunca permitiu que eu questionasse meu potencial e esteve ao meu lado em todos os momentos: levarei para a vida todas as boas (e as não tão boas) práticas que aprendi com vocês, que servem de motivação para o professor e pesquisador que almejo me tornar. Obrigado pela confiança e pela oportunidade de trabalhar com vocês. O trabalho não acaba aqui.

“Se eu soubesse antes o que sei agora, erraria tudo exatamente igual.”

“If I knew then what I know now, I would have made exactly the same mistakes.”

From *Surfando Karmas & DNA* (2001) by Engenheiros do Hawaii

“So, nobody told me what to do, and there was no preconception of what to do.”

From *Giorgio By Moroder* (2013) by Daft Punk

Abstract

In statistical modelling with historical data, a central challenge is how to effectively integrate prior evidence with current observations. This issue is particularly pressing in clinical trials and medical applications, where using historical data can improve both the efficiency and robustness of inference. A natural Bayesian strategy for this purpose is the use of informative priors. Among them, the normalised power prior has attracted significant interest, as it provides a principled way to control the influence of historical information. The method introduces a power parameter, $\eta \in [0, 1]$, that governs the degree of borrowing: $\eta = 0$ discards the past data while $\eta = 1$ fully incorporates it. This flexibility makes the framework highly appealing in practice. Despite its advantages, the normalised power prior induces a doubly intractable posterior: the prior itself contains a normalising constant $c_0(\eta)$ that depends on the power parameter η and is unavailable in closed form for most models of interest, rendering standard Metropolis–Hastings inapplicable. Existing approaches therefore rely on approximate inference methods that often lack theoretical guarantees and explicit convergence properties. To overcome these limitations, we propose exact Markov chain Monte Carlo algorithms that avoid direct evaluation of the normalising constant ratio. The key insight is a thermodynamic identity expressing the log-ratio of normalising constants as a tractable conditional expectation under the power prior. Exploiting this identity, we construct non-negative unbiased estimators of the normalising constant ratio via Monte Carlo integration, which we embed within an Exchange Metropolis–Hastings scheme and a pseudo-marginal framework. We evaluate the method on Gaussian models under varying degrees of compatibility between historical and current data. Results demonstrate that the algorithm accurately reflects data conflict by shrinking η when the datasets are incompatible, while recovering correct posterior behaviour when they are aligned. Comparisons with analytical baselines confirm that our method achieves reliable estimates of both model parameters and the power parameter, with efficiency suitable for practical use. This provides a rigorous, exact alternative to approximate approaches, offering stronger guarantees for applications such as adaptive clinical trial design.

Keywords: Bayesian Inference; Power Prior; Doubly Intractable Distributions; Exact MCMC; Historical Data Analysis

Resumo

Na modelagem estatística com dados históricos, um desafio central é integrar efetivamente evidências prévias com observações correntes. Essa questão é especialmente relevante em ensaios clínicos e aplicações médicas, nos quais o uso de dados históricos pode aumentar tanto a eficiência quanto a robustez da inferência. Uma estratégia bayesiana natural para esse fim é o uso de prioris informativas. Dentre elas, a Normalised Power Prior tem atraído interesse considerável por oferecer uma forma principiada de controlar a influência das informações históricas. O método introduz um parâmetro de potência, $\eta \in [0, 1]$, que governa o grau de incorporação: $\eta = 0$ descarta os dados passados, ao passo que $\eta = 1$ os incorpora integralmente. Essa flexibilidade torna o arcabouço bastante atraente na prática. Apesar de suas vantagens, a Normalised Power Prior induz uma posteriori duplamente intratável: a priori em si contém uma constante normalizadora $c_0(\eta)$ que depende do parâmetro de potência η e não está disponível em forma fechada para a maioria dos modelos de interesse, inviabilizando a aplicação direta do algoritmo Metropolis–Hastings padrão. As abordagens existentes dependem, portanto, de métodos de inferência aproximada que frequentemente carecem de garantias teóricas e propriedades de convergência explícitas. Para superar essas limitações, propomos algoritmos exatos de Markov chain Monte Carlo que evitam a avaliação direta da razão de constantes de normalização. A ideia central é uma identidade termodinâmica que expressa a log-razão das constantes de normalização como uma esperança condicional tratável sob a power prior. Explorando essa identidade, construímos estimadores não negativos e não viesados da razão de constantes de normalização via integração de Monte Carlo, os quais são incorporados em um esquema Exchange Metropolis–Hastings e em um método pseudo-marginal. Avaliamos o método em modelos gaussianos sob diferentes graus de compatibilidade entre dados históricos e correntes. Os resultados demonstram que o algoritmo reflete adequadamente o conflito entre os dados ao encolher η quando os conjuntos são incompatíveis, ao mesmo tempo que recupera o comportamento *a posteriori* correto quando estão alinhados. Comparações com referenciais analíticos confirmam que o método obtém estimativas confiáveis tanto dos parâmetros do modelo quanto do parâmetro de potência, com eficiência adequada ao uso prático. Isso fornece uma alternativa exata e rigorosa às abordagens aproximadas, oferecendo garantias mais sólidas para aplicações como o planejamento adaptativo de ensaios clínicos.

Palavras-chave: Inferência Bayesiana; Power Prior; Distribuições Duplamente Intratáveis; MCMC Exato; Análise de Dados Históricos

List of Figures

Figure 1 – Posterior density of η across discrepancy regimes for the Gaussian normalised power prior experiment.	48
Figure 2 – Posterior density of $\ \theta\ _2$ across discrepancy regimes for the Gaussian normalised power prior experiment.	49
Figure 3 – Posterior density of η for the linear regression NPP experiment across discrepancy scenarios.	52
Figure 4 – Posterior density of $\ \beta\ _2$ for the linear regression NPP experiment across discrepancy scenarios.	53
Figure 5 – Posterior density of σ^2 for the linear regression NPP experiment across discrepancy scenarios.	53
Figure 6 – Posterior density of η across discrepancy scenarios for the logistic regression normalised power prior experiment.	56
Figure 7 – Posterior density of the first regression coefficient across discrepancy scenarios.	57
Figure 8 – Posterior density of the second regression coefficient across discrepancy scenarios.	57
Figure 9 – Posterior density of the third regression coefficient across discrepancy scenarios.	57

List of Tables

Table 1	– Experiment 1 convergence diagnostics. Batches: number of batches of 4 chains $\times$ 5000 iterations required to reach $W_1(\eta) < 0.01$. Wasserstein distances computed against the analytic MH reference.	49
Table 2	– Experiment 1 posterior summaries (Normal-Normal NPP).	50
Table 3	– Diagnostic comparison for the linear regression NPP experiment with Normal-Inverse-Gamma priors. Wasserstein distances are computed against the analytic MH baseline within each discrepancy scenario. . . .	53
Table 4	– Posterior summaries for the NIG linear regression NPP experiment. . .	54
Table 5	– Diagnostic comparison for the logistic regression NPP experiment. Wasserstein distances are computed against the exchange MH reference within each discrepancy scenario.	58
Table 6	– Posterior summaries for the logistic regression NPP experiment across three discrepancy levels.	58

Contents

1	INTRODUCTION	12
2	BACKGROUND	13
2.1	Bayesian inference	13
2.1.1	Principles	13
2.1.2	Monte Carlo methods	14
2.2	Classical Markov chain Monte Carlo methods	15
2.2.1	Gibbs sampling	16
2.2.2	Metropolis–Hastings	17
2.3	Doubly intractable distributions	18
2.4	The Normalised Power Prior	19
2.4.1	Formulation	20
2.4.2	The derivative of the normalising constant	21
2.4.3	Other historical data modelling approaches	21
2.5	MCMC methods for intractable normalising constants	21
2.5.1	The exchange algorithm	22
2.5.2	Barker’s algorithm	23
2.5.2.1	The two-coin algorithm	23
2.5.3	Divide-and-conquer Bernoulli factory	24
2.5.3.1	The coin merge procedure	25
2.5.3.2	Hierarchical factorization	25
2.5.3.3	Computational complexity	26
2.5.4	Pseudo-marginal Metropolis–Hastings	27
2.6	Function emulation methods	28
2.6.1	Gaussian processes	29
2.6.2	Generalised additive models	29
3	APPROXIMATE INFERENCE WITH NOISY-MCMC	31
3.1	Emulating the normalising constant	31
3.1.1	Consequences	31
3.1.2	Shape-constrained Gaussian process emulation	32
3.1.3	Shape-constrained additive models	33
3.2	Current approaches for the NPP	33
3.2.1	Approximate inference through emulation and interpolation	34
3.2.2	Monte Carlo approaches and exact inference	35
3.3	Function emulation for doubly intractable settings	36

4	EXACT INFERENCE FOR THE NORMALISED POWER PRIOR	38
4.1	Exchange Metropolis–Hastings	38
4.2	Barker’s method via Poisson estimator	39
4.2.1	Producing an unbiased estimator via a Poisson estimator	39
4.2.1.1	Practical implementation	41
4.2.2	Local envelopes for the Poisson estimator	41
4.2.3	Divide-and-conquer factorization	42
4.3	Pseudo-marginal via importance sampling	44
5	EXPERIMENTS	46
5.1	Assessing exact and approximate samplers	46
5.1.1	Using exact methods as baseline	46
5.2	Convergence to the baseline	47
5.2.1	Normal distribution with known covariance	47
5.2.1.1	Discussion of results	49
5.2.2	Linear regression with Normal-Inverse-Gamma priors	50
5.2.2.1	Discussion of results	54
5.2.3	Logistic regression with surrogate normalising constants	55
5.2.3.1	Discussion of results	58
6	DISCUSSION	60
6.1	Summary of contributions	60
6.2	Empirical findings	60
6.3	Limitations and future directions	61
	REFERENCES	62
	APPENDIX A – PROOFS	65
A.1	Thermodynamic identity for the normalising constant	65
A.2	Unbiasedness of the Poisson estimator	65
A.3	Correctness of the exchange algorithm	66
A.4	Convexity of log-normalising constant	66
	APPENDIX B – ANALYTIC FORMULATIONS	67
B.1	Gaussian likelihood with known variance	67
B.1.1	Adding a Gaussian initial prior	68
B.2	Linear regression with NIG priors	70

1 Introduction

In statistical modelling with historical data, a central challenge is effectively integrating information from past studies with current data. This is particularly relevant in clinical trials and medical research, where leveraging historical data can improve the efficiency and robustness of statistical inferences. A natural approach in this setting is the use of *informative priors*, which incorporate prior knowledge into the analysis, thereby enhancing parameter estimation and decision-making (Spiegelhalter, 2004; Gelman et al., 2013; Ibrahim et al., 2015).

Among the various methods for constructing informative priors, *power priors* (Ibrahim; Chen, 2000; Ibrahim et al., 2015) have gained significant attention. They allow for controlled borrowing of information from historical data by adjusting a power parameter, η , which typically ranges between 0 and 1. When $\eta = 0$, no historical information is incorporated, whereas $\eta = 1$ fully integrates the historical data into the current analysis. This flexibility makes power priors particularly appealing in adaptive clinical trial design and other settings where historical evidence is available, but its relevance is uncertain.

However, when η is itself treated as unknown and assigned a prior distribution, the resulting posterior is *doubly intractable*: the normalised power prior contains a normalising constant $c_0(\eta)$ that depends on η and cannot be evaluated in closed form for general likelihoods, preventing direct use of Metropolis–Hastings. Current approaches rely on *approximate methods* (Carvalho; Ibrahim, 2021) that lack theoretical guarantees and explicit convergence bounds. This work addresses this gap by developing exact Markov chain Monte Carlo algorithms (Murray; Ghahramani; MacKay, 2012; Gonçalves; Łatuszyński; Roberts, 2017; Vats et al., 2021) that avoid evaluating the intractable normalising constant ratio altogether, providing a rigorous computational baseline for inference under the normalised power prior.

2 Background

2.1 Bayesian inference

2.1.1 Principles

Bayesian inference provides a probabilistic framework for updating uncertainty about unknown quantities after observing data.

Let $(\Theta, \mathcal{F})$ denote the parameter space and let D denote the observed data. Assume that the statistical model $\{P_\theta : \theta \in \Theta\}$ is dominated by a common measure μ , with likelihood $L(D | \theta)$. Let P_0 denote a prior distribution on $(\Theta, \mathcal{F})$. Whenever P_0 admits a density π_0 with respect to a reference measure ν on Θ , we write

$$P_0(d\theta) = \pi_0(\theta) \nu(d\theta). \quad (2.1)$$

The Bayesian model is then specified by the likelihood function $L(D | \theta)$ and the prior density $\pi_0(\theta)$. The likelihood describes the information about θ contained in the data, while the prior represents information available before observing D . Bayes' theorem combines these two sources of information through the posterior distribution

$$\pi(\theta | D) = \frac{L(D | \theta) \pi_0(\theta)}{\int_{\Theta} L(D | t) \pi_0(t) dt}. \quad (2.2)$$

The denominator in (2.2) is the marginal likelihood, or evidence, and is given by

$$m(D) = \int_{\Theta} L(D | t) \pi_0(t) dt. \quad (2.3)$$

When $m(D) < \infty$, the posterior distribution is well-defined and integrates to one. In this sense, Bayes' theorem updates the prior distribution through the likelihood, with the marginal likelihood ensuring that the resulting posterior is a proper probability distribution on Θ .

For any posterior-integrable function $h : \Theta \rightarrow \mathbb{R}$, posterior summaries are often expressed through expectations of the form

$$\mathbb{E}_{\pi(\theta|D)}[h(\theta)] = \int_{\Theta} h(\theta) \pi(\theta | D) d\theta. \quad (2.4)$$

Examples include posterior means, posterior variances, marginal probabilities, credible intervals, and posterior predictive quantities. The main computational challenge is that the posterior distribution is usually available only up to a normalising constant.

Indeed, many Bayesian algorithms only require the unnormalised posterior density

$$\gamma(\theta) = L(D \mid \theta)\pi_0(\theta), \quad (2.5)$$

because the marginal likelihood $m(D)$ does not depend on θ .

However, the marginal likelihood is still important for model comparison, empirical Bayes procedures, and doubly intractable models. In low-dimensional conjugate models, the integral in (2.3) may be available analytically.

In modern Bayesian models, this integral is often unavailable, and posterior expectations such as (2.4) must be approximated numerically. This motivates the use of simulation-based methods.

2.1.2 Monte Carlo methods

Monte Carlo methods approximate integrals by random sampling.

Let π be a probability distribution on $(\Theta, \mathcal{F})$, and suppose that the quantity of interest is

$$\mu = \mathbb{E}_\pi[h(\theta)] = \int_\Theta h(\theta)\pi(\theta) \, d\theta. \quad (2.6)$$

If $\theta_1, \dots, \theta_N$ are independent samples from π , then the Monte Carlo estimator of μ is

$$\hat{\mu}_N = \frac{1}{N} \sum_{i=1}^N h(\theta_i). \quad (2.7)$$

Under standard integrability conditions, the law of large numbers gives

$$\hat{\mu}_N \xrightarrow[N \rightarrow \infty]{a.s.} \mu. \quad (2.8)$$

If $h(\theta)$ has finite variance under π , the central limit theorem gives

$$\sqrt{N}(\hat{\mu}_N - \mu) \xrightarrow[N \rightarrow \infty]{d} \mathcal{N}(0, \sigma_h^2), \quad (2.9)$$

where

$$\sigma_h^2 = \text{Var}_\pi [h(\theta)]. \quad (2.10)$$

The important feature of the Monte Carlo approximation is that its convergence rate is typically $N^{-1/2}$ and does not directly depend on the dimension of Θ . This makes Monte Carlo methods attractive for Bayesian inference, where posterior distributions are often high-dimensional.

The difficulty is that independent sampling from $\pi(\theta \mid D)$ is rarely possible. Markov chain Monte Carlo methods address this difficulty by constructing a Markov chain whose invariant distribution is the desired posterior distribution.

The samples are then dependent, but posterior expectations can still be estimated using ergodic averages. If $(\theta_n)_{n \geq 0}$ is a Markov chain with invariant distribution π , then the MCMC estimator is

$$\hat{\mu}_N^{\text{MCMC}} = \frac{1}{N} \sum_{n=1}^N h(\theta_n). \quad (2.11)$$

Under suitable irreducibility, recurrence, and integrability conditions, the ergodic theorem gives ([Robert; Casella, 2004](#); [Meyn; Tweedie; Glynn, 2009](#); [Polasek, 2012](#))

$$\hat{\mu}_N^{\text{MCMC}} \xrightarrow[N \rightarrow \infty]{a.s.} \mathbb{E}_\pi [h(\theta)]. \quad (2.12)$$

The price paid for replacing independent samples with a Markov chain is autocorrelation.

When samples are correlated, the effective amount of information in N iterations may be much smaller than N independent draws. This motivates the use of diagnostics such as effective sample size, Monte Carlo standard errors, trace plots, and comparisons across independent chains. Another important practical issue is initialization.

A Markov chain usually starts from a distribution different from π , and therefore the early iterations may not be representative of the target distribution. The usual practical response is to discard an initial portion of the chain, known as burn-in.

However, burn-in is an asymptotic correction rather than an exact one. This distinction is central to the later discussion of exact and unbiased MCMC methods.

2.2 Classical Markov chain Monte Carlo methods

Markov chain Monte Carlo methods construct a transition kernel P on $(\Theta, \mathcal{F})$ such that the posterior distribution π is invariant.

This means that, for every measurable set $A \in \mathcal{F}$,

$$\int_{\Theta} P(\theta, A) \pi(\theta) d\theta = \pi(A). \quad (2.13)$$

Equivalently, if $\theta_n \sim \pi$ and $\theta_{n+1} \mid \theta_n \sim P(\theta_n, \cdot)$, then $\theta_{n+1} \sim \pi$. Invariance alone does not guarantee convergence from arbitrary starting points.

For convergence, the Markov chain must also satisfy appropriate irreducibility and recurrence conditions. A common sufficient condition used to establish invariance is detailed balance (Meyn; Tweedie; Glynn, 2009).

The transition kernel P satisfies detailed balance with respect to π if

$$\pi(\theta) P(\theta, d\theta') = \pi(\theta') P(\theta', d\theta) \quad (2.14)$$

as measures on $\Theta \times \Theta$.

Detailed balance implies invariance and gives a convenient route for proving that a Markov chain targets the desired posterior distribution. The two classical algorithms discussed in this section are Gibbs sampling and Metropolis–Hastings. Gibbs sampling exploits conditional distributions that can be exactly sampled from.

Metropolis–Hastings is more general and only requires evaluation of the target density up to a normalising constant. Both algorithms are asymptotically exact in the sense that their ergodic averages converge to posterior expectations under suitable conditions. They are not exact in the finite-time sense, because samples obtained after a finite number of iterations still depend on the initial state.

2.2.1 Gibbs sampling

Gibbs sampling (Geman; Geman, 1984; Gelfand; Smith, 1990) is designed for posterior distributions whose full conditional distributions are available and easy to sample from.

Suppose that the parameter can be partitioned into K blocks,

$$\theta = (\theta_1, \dots, \theta_K), \quad \Theta = \Theta_1 \times \dots \times \Theta_K. \quad (2.15)$$

For each block θ_k , define the complementary vector

$$\theta_{-k} = (\theta_1, \dots, \theta_{k-1}, \theta_{k+1}, \dots, \theta_K). \quad (2.16)$$

The full conditional distribution of the k -th block is

$$\pi(\theta_k \mid \theta_{-k}, D). \quad (2.17)$$

A systematic-scan Gibbs sampler updates the blocks sequentially. Starting from $\theta^{(n)} = (\theta_1^{(n)}, \dots, \theta_K^{(n)})$, one iteration produces $\theta^{(n+1)}$ by sampling

$$\theta_1^{(n+1)} \sim \pi(\theta_1 \mid \theta_2^{(n)}, \dots, \theta_K^{(n)}, D), \quad (2.18)$$

$$\theta_2^{(n+1)} \sim \pi(\theta_2 \mid \theta_1^{(n+1)}, \theta_3^{(n)}, \dots, \theta_K^{(n)}, D), \quad (2.19)$$

$$\vdots$$

$$\theta_K^{(n+1)} \sim \pi(\theta_K \mid \theta_1^{(n+1)}, \dots, \theta_{K-1}^{(n+1)}, D). \quad (2.20)$$

Each update leaves the posterior distribution invariant because it samples exactly from a conditional distribution of the target. The composition of these invariant updates also leaves the posterior distribution invariant. The main advantage of Gibbs sampling is that every proposed block update is accepted automatically. In this sense, the algorithm avoids the tuning of an acceptance probability.

On the other hand, the most relevant limitation is that the required full conditional distributions may not be available in closed form. Even when they are available, strong posterior dependence between blocks can lead to slow mixing. This is especially relevant in hierarchical models, latent variable models, and models with highly correlated parameters. In such cases, blocking, reparameterization, or hybrid Gibbs-within-Metropolis strategies may be necessary (Gilks; Best; Tan, 1995).

2.2.2 Metropolis–Hastings

The Metropolis–Hastings (Metropolis et al., 1953; Hastings, 1970) algorithm constructs a Markov chain using a proposal distribution and an accept-reject correction. Let π denote the target posterior density, known up to a normalising constant.

Let $q(\theta' \mid \theta)$ be a proposal density for moving from the current state θ to a proposed state θ' . Given the current state θ , the algorithm first samples

$$\theta' \sim q(\cdot \mid \theta). \quad (2.21)$$

The proposal is then accepted with probability

$$\alpha(\theta, \theta') = \min \left\{ 1, \frac{\pi(\theta')q(\theta \mid \theta')}{\pi(\theta)q(\theta' \mid \theta)} \right\}. \quad (2.22)$$

If the proposal is accepted, the next state is θ' . If the proposal is rejected, the chain remains at θ . Thus,

$$\theta_{n+1} = \begin{cases} \theta', & \text{with probability } \alpha(\theta_n, \theta'), \\ \theta_n, & \text{otherwise.} \end{cases} \quad (2.23)$$

The acceptance probability in (2.22) depends only on ratios of the target density. Therefore, if $\pi(\theta) = \gamma(\theta)/Z$ for an unknown normalising constant Z , then Z cancels from the ratio. This cancellation is one of the main reasons why Metropolis–Hastings is useful in Bayesian computation.

The Metropolis–Hastings transition kernel satisfies detailed balance with respect to π . Indeed, the accept-reject correction symmetrizes the probability flow between θ and θ' . Consequently, π is invariant for the Markov chain. A common special case is the random-walk Metropolis algorithm. In this case,

$$q(\theta' \mid \theta) = q(\theta - \theta'), \quad (2.24)$$

so that $q(\theta' \mid \theta) = q(\theta \mid \theta')$. The acceptance probability then reduces to

$$\alpha(\theta, \theta') = \min \left\{ 1, \frac{\pi(\theta')}{\pi(\theta)} \right\}. \quad (2.25)$$

The efficiency of Metropolis–Hastings depends strongly on the proposal distribution. If the proposal variance is too small, the chain accepts many proposals but explores the state space slowly. If the proposal variance is too large, the chain proposes distant states that are often rejected. This trade-off is one of the central practical issues in MCMC.

Metropolis–Hastings is asymptotically exact under appropriate conditions, but it does not remove the finite-time dependence on the initial state. For this reason, convergence diagnostics and burn-in choices remain important in practical applications. In doubly intractable models, even the ratio in (2.22) may contain unknown normalising constants that depend on θ . This is one of the settings where standard Metropolis–Hastings is no longer directly applicable and where the methods studied in this dissertation become necessary.

2.3 Doubly intractable distributions

A posterior distribution is called *doubly intractable* when its unnormalised density cannot be computed pointwise (Murray; Ghahramani; MacKay, 2012; Park; Haran, 2020). The difficulty is distinct from the ordinary intractability of the marginal likelihood $m(D)$,

which cancels from the Metropolis–Hastings ratio and therefore poses no algorithmic obstacle. In a doubly intractable model, the posterior density takes the form

$$\pi(\theta \mid D) \propto \frac{\pi_0(\theta) h(D \mid \theta)}{Z(\theta)}, \quad (2.26)$$

where $Z(\theta)$ depends on the parameter θ and cannot be evaluated. The ratio $Z(\theta)/Z(\theta')$ required by the Metropolis–Hastings acceptance step is therefore itself intractable: both the numerator and the denominator involve the same unknown normalising function evaluated at different parameter values. This is the second source of intractability, in addition to the marginal likelihood.

Doubly intractable models arise naturally in several important classes. Exponential random graph models have a likelihood of the form $\exp\{\theta^\top t(y)\}/Z(\theta)$, where $Z(\theta)$ sums over all graphs of a given size and is computationally infeasible for all but very small networks. Gibbs random field models and Markov point processes share the same structure.

The normalised power prior, introduced in Section 2.4, belongs to this class for a different reason: here $Z(\theta)$ is not a partition function of the likelihood but the normalising constant $c_0(\eta)$ of the *prior* distribution, which depends on the power parameter η . This is an unusual member of the doubly intractable family, in that the intractable constant arises through the Bayesian modelling choice rather than through the likelihood structure itself. The methods reviewed in Section 2.5 and those developed in Chapter 4 are designed to handle precisely this setting.

2.4 The Normalised Power Prior

In many applied settings, a new study is conducted in a context where data from previous, related studies are already available. Clinical trials are a canonical example: regulatory agencies and trial designers routinely face the question of how much weight to give historical control data when planning or analysing a new trial (Spiegelhalter, 2004; Ibrahim et al., 2015). A natural Bayesian strategy is to use the historical data to construct an informative prior for the current analysis. However, doing so naively, i.e., by treating the historical data as part of the current dataset, ignores the possibility that the two data sources are not fully exchangeable, for instance due to changes in patient populations, measurement protocols, or underlying disease dynamics.

The power prior framework of Chen and Ibrahim (2000) addresses this by introducing a scalar parameter $\eta \in [0, 1]$ that controls how much influence the historical data exert on the prior. Rather than fully incorporating historical evidence, the prior is formed by raising the historical likelihood to the power η before combining it with an initial prior. When $\eta = 0$, the historical data are ignored entirely; when $\eta = 1$, they are weighted equally to a single observation from the current study. Intermediate values allow a continuum of

borrowing, and treating η as unknown and inferring it from the data provides automatic, data-driven borrowing strength.

The normalised power prior (Duan; Ye; Smith, 2006; Neuenschwander; Branson; Spiegelhalter, 2009; Carvalho; Ibrahim, 2021) is the version of the power prior that properly accounts for the normalising constant of the powered likelihood, ensuring that η has a coherent joint prior interpretation alongside θ . As established in Section 2.3, this normalisation introduces a parameter-dependent constant $c_0(\eta)$ that renders the posterior doubly intractable. This computational difficulty is the central challenge that this thesis addresses.

2.4.1 Formulation

Recalling the framework of Section 2.1, the posterior distribution of θ given data D takes the form of (2.2), with $\pi_0(\theta)$ serving as the initial prior. The normalised power prior framework replaces this initial prior by a data-informed version constructed from historical data $D_0 = \{d_{01}, \dots, d_{0N_0}\}$.

Definition 2.4.1 (Power prior). Let $D_0 = \{d_{01}, d_{02}, \dots, d_{0N_0}\}$ with $d_{0i} \in \mathcal{X}^p$ be the historical data, and let $L_0(D_0 | \theta)$ be a likelihood function that is finite for all $\theta \in \Theta \subseteq \mathbb{R}^p$. Then, the power prior is defined as

$$\pi_\eta(\theta | D_0) \propto L_0(D_0 | \theta)^\eta \pi_0(\theta), \quad (2.27)$$

where $\pi_0(\theta)$ is the initial prior distribution of θ and η is the (fixed) power parameter.

To accommodate uncertainty in the choice of η , we treat it as a random variable and assign a prior distribution $\pi_A(\eta | \varphi)$, indexed by hyperparameters φ .

The unnormalised joint power prior over (θ, η) is then

$$\pi^*(\theta, \eta | D_0, \varphi) \propto L_0(D_0 | \theta)^\eta \pi_0(\theta) \pi_A(\eta | \varphi). \quad (2.28)$$

To normalise this joint prior, one requires the normalising constant

$$c_0(\eta) = \int_{\Theta} L_0(D_0 | t)^\eta dP_0(t), \quad (2.29)$$

which is unavailable in closed form for general likelihoods.

Given new data D , the marginal posterior distribution of η is

$$\pi(\eta | D, D_0, \varphi) \propto \frac{\pi_A(\eta | \varphi)}{c_0(\eta)} \int_{\Theta} L(D | t) L_0(D_0 | t)^\eta dP_0(t). \quad (2.30)$$

Since neither $c_0(\eta)$ nor the integral in (2.30) is tractable in general, this is a *doubly intractable* posterior in the sense of Section 2.3. The computational challenge this poses, and the strategies developed to address it, are the central subject of this thesis.

2.4.2 The derivative of the normalising constant

Under regularity conditions permitting differentiation under the integral sign, the thermodynamic identity (proved in Appendix A.1 and discussed by Friel and Pettitt (2008)) gives

$$\frac{d \log c_0(\eta)}{d\eta} = \mathbb{E}_{\pi_\eta}[\log L_0(D_0 \mid \theta)], \quad (2.31)$$

where $\mathbb{E}_{\pi_\eta}[\cdot]$ denotes expectation under the power prior $\pi_\eta(\theta \mid D_0)$.

This identity plays a central role in the exact inference methods of Chapter 4: it expresses the log-derivative of an intractable normalising constant as a tractable posterior expectation, which can be estimated by Monte Carlo simulation.

2.4.3 Other historical data modelling approaches

The normalised power prior is one among several principled Bayesian approaches to incorporating historical information. The *commensurate prior* of Hobbs et al. (2011) models the relationship between historical and current parameters through a shared hierarchical structure, allowing the degree of borrowing to adapt to the compatibility between studies.

Meta-analytic-predictive (MAP) priors (Neuenschwander et al., 2010) derive an informative prior for the current study by fitting a hierarchical model to a collection of historical studies and marginalising over the between-study heterogeneity.

Robust mixture priors (Schmidli et al., 2014) address the risk of prior-data conflict by mixing an informative component derived from historical evidence with a vague component, ensuring that the prior retains a non-negligible weight in regions not supported by the historical data.

Each of these alternatives makes different modelling assumptions about the exchangeability of the historical and current parameters, and each leads to different computational challenges. The normalised power prior is particularly appealing when a single scalar parameter $\eta \in [0, 1]$ suffices to capture the degree of historical relevance, and when the historical and current data share the same likelihood structure.

2.5 MCMC methods for intractable normalising constants

Many Bayesian models involve posterior distributions whose density can only be evaluated up to a normalising constant. This is not an obstacle for ordinary Metropolis–Hastings when the unknown normalising constant does not depend on the parameter being updated. However, the situation is considerably harder when the normalising constant itself depends on the parameter, as is the case for the normalised power prior.

In this section, we review three families of methods that are relevant for this thesis. First, we discuss the exchange algorithm, which removes the intractable ratio of normalising constants by introducing an auxiliary draw from a proposal-dependent distribution. Second, we discuss Barker’s algorithm and Bernoulli factory methods, which replace the explicit evaluation of an acceptance probability by an exact randomisation procedure. Third, we discuss divide-and-conquer Bernoulli factories, which aim to reduce the computational cost of Bernoulli factory methods in models with many likelihood factors.

2.5.1 The exchange algorithm

Consider a family of unnormalised densities $f_\eta(\theta)$ with normalising constant

$$c(\eta) = \int_{\Theta} f_\eta(\theta) d\theta. \quad (2.32)$$

Suppose that the target distribution has density

$$\pi(\eta \mid y) \propto \frac{\gamma(\eta \mid y)}{c(\eta)}. \quad (2.33)$$

A standard Metropolis–Hastings proposal, as discussed in Section 2.2.2, from η to η' requires the ratio

$$\frac{\pi(\eta' \mid y)}{\pi(\eta \mid y)} = \frac{\gamma(\eta' \mid y)}{\gamma(\eta \mid y)} \frac{c(\eta)}{c(\eta')}. \quad (2.34)$$

The ratio $c(\eta)/c(\eta')$ is generally unavailable. This prevents the direct application of ordinary Metropolis–Hastings.

The exchange algorithm of [Murray, Ghahramani and MacKay \(2012\)](#) introduces an auxiliary variable $x^\star$ sampled exactly from the normalised model at the proposed value η' . That is,

$$x^\star \sim \frac{f_{\eta'}(x)}{c(\eta')}. \quad (2.35)$$

The Metropolis–Hastings ratio is then constructed on an augmented space so that the unknown normalising constants cancel. In its simplest form, the exchange ratio is

$$R_{\text{ex}}(\eta, \eta'; x^\star) = \frac{\gamma(\eta' \mid y)}{\gamma(\eta \mid y)} \frac{q(\eta \mid \eta')}{q(\eta' \mid \eta)} \frac{f_\eta(x^\star)}{f_{\eta'}(x^\star)}. \quad (2.36)$$

The corresponding acceptance probability is

$$\alpha_{\text{ex}}(\eta, \eta') = \min \{1, R_{\text{ex}}(\eta, \eta'; x^\star)\}. \quad (2.37)$$

The key point is that R_{ex} contains only unnormalised densities and proposal terms. Therefore, the intractable ratio $c(\eta)/c(\eta')$ is not evaluated directly.

The exchange algorithm is particularly well-suited to settings where exact simulation from the normalised model is available.

2.5.2 Barker's algorithm

Barker's algorithm (Barker, 1965) is an alternative accept–reject rule for Markov chain Monte Carlo. For a target distribution $\pi(\theta)$ and proposal density $q(\theta' | \theta)$, define the usual Metropolis–Hastings odds ratio as

$$R(\theta, \theta') = \frac{\pi(\theta')q(\theta | \theta')}{\pi(\theta)q(\theta' | \theta)}. \quad (2.38)$$

The Barker acceptance probability is

$$\alpha_B(\theta, \theta') = \frac{R(\theta, \theta')}{1 + R(\theta, \theta')}. \quad (2.39)$$

Equivalently,

$$\alpha_B(\theta, \theta') = \frac{\pi(\theta')q(\theta | \theta')}{\pi(\theta)q(\theta' | \theta) + \pi(\theta')q(\theta | \theta')}. \quad (2.40)$$

This acceptance probability satisfies detailed balance because

$$\frac{\alpha_B(\theta, \theta')}{\alpha_B(\theta', \theta)} = R(\theta, \theta'). \quad (2.41)$$

Therefore, Barker's algorithm leaves the target distribution invariant.

Barker's algorithm is usually less efficient than Metropolis–Hastings in the Peskun ordering (Gonçalves; Łatuszyński; Roberts, 2017). However, its acceptance probability has a useful algebraic form. In particular, it can be implemented by Bernoulli factory methods when the odds ratio can be represented through probabilities that can be simulated but not evaluated. This makes Barker's algorithm useful for some intractable likelihood problems (Gonçalves; Łatuszyński; Roberts, 2017; Vats et al., 2021).

The important distinction is that Barker's algorithm is not required for the exchange algorithm. The exchange algorithm can be combined with the ordinary Metropolis–Hastings acceptance probability. Barker's rule becomes useful when the goal is to implement an accept–reject step through Bernoulli factories rather than through direct numerical evaluation of the acceptance probability.

2.5.2.1 The two-coin algorithm

The two-coin algorithm is a basic Bernoulli factory construction used to implement Barker-type accept–reject probabilities (Gonçalves; Łatuszyński; Roberts, 2017). Suppose that the odds ratio can be factorized as

$$R(\theta, \theta') = \frac{c(\theta', \theta)p(\theta', \theta)}{c(\theta, \theta')p(\theta, \theta')}, \quad (2.42)$$

where $c(\cdot, \cdot)$ is tractable and $p(\cdot, \cdot)$ is a probability that can be simulated. Then Barker's acceptance probability can be written as

$$\alpha_B(\theta, \theta') = \frac{c(\theta', \theta)p(\theta', \theta)}{c(\theta', \theta)p(\theta', \theta) + c(\theta, \theta')p(\theta, \theta')}. \quad (2.43)$$

The two-coin algorithm generates a Bernoulli random variable with this probability without evaluating $p(\theta', \theta)$ or $p(\theta, \theta')$ numerically.

Algorithm 1: Two-coin algorithm

```

1 Flip a coin with probability proportional to  $c(\theta', \theta)$  against  $c(\theta, \theta')$ ;
2 if the first-stage coin selects the proposed direction then
3   | Flip a coin with probability  $p(\theta', \theta)$ ;
4   | if the second-stage coin is heads then
5   |   | accept the proposal;
6   | end
7 else
8   | Flip a coin with probability  $p(\theta, \theta')$ ;
9   | if the second-stage coin is heads then
10  |   | reject the proposal;
11  | end
12 if no decision has been made then
13  | return to the first step;
14 end

```

The algorithm loops until one of the simulated coins produces a decisive event. Conditioning on the first decisive event gives the Barker probability. Thus, the resulting accept–reject decision is exact even when the probabilities $p(\theta', \theta)$ and $p(\theta, \theta')$ are only available through simulation.

The practical difficulty is that the expected number of loops can be large. This happens when the decisive events are rare. In models with many likelihood factors, this cost can grow rapidly with the data size.

2.5.3 Divide-and-conquer Bernoulli factory

For models with independent observations, the likelihood often factorizes as

$$L(D \mid \theta) = \prod_{i=1}^n L_i(\theta). \quad (2.44)$$

A Bernoulli factory implementation may then require simulating a coin whose probability is a product of many factor probabilities. If the required probability is

$$p(\theta) = \prod_{i=1}^n p_i(\theta), \quad (2.45)$$

then $p(\theta)$ can become very small as n increases. A direct two-coin implementation may therefore have an expected number of loops that grows exponentially in n . This is the main computational bottleneck addressed by divide-and-conquer Bernoulli factories.

The divide-and-conquer Bernoulli factory of [Stumpf-Fétizon and Gonçalves \(2025\)](#) reduces this cost by recursively merging coins associated with smaller groups of factors. Instead of constructing one coin for the full product at once, the method builds a tree. The leaves correspond to individual factors or small batches of factors. The internal nodes merge the coins produced by their children. The root returns the final Bernoulli variable required by the Barker accept–reject step.

2.5.3.1 The coin merge procedure

The basic merge identity can be described in terms of odds. Suppose that two independent coins have success probabilities

$$r_1 = \frac{h_1}{h_1 + g_1} \quad (2.46)$$

and

$$r_2 = \frac{h_2}{h_2 + g_2}. \quad (2.47)$$

The goal is to construct a coin with success probability

$$r_0 = \frac{h_1 h_2}{h_1 h_2 + g_1 g_2}. \quad (2.48)$$

The merge procedure repeatedly flips the two child coins until they agree. If both child coins return one, the merged coin returns one. If both child coins return zero, the merged coin returns zero. If the two child coins disagree, the procedure is repeated.

Algorithm 2: Coin merge algorithm

```

1 Flip independent coins with probabilities  $r_1$  and  $r_2$ ;
2 if both coins return the same value then
3   | return the common value;
4 else
5   | return to the first step;
6 end
```

The probability that the merged coin returns one is

$$\mathbb{P}(\text{return 1}) = \frac{r_1 r_2}{r_1 r_2 + (1 - r_1)(1 - r_2)} \quad (2.49)$$

$$= \frac{h_1 h_2}{h_1 h_2 + g_1 g_2}. \quad (2.50)$$

Thus, the merge operation multiplies the odds represented by the two child coins.

2.5.3.2 Hierarchical factorization

Given n factors, the DCBF organizes them in a binary tree. At the leaves, each node represents either one likelihood factor or a small batch of likelihood factors. At each

internal node, the two child coins are merged using Algorithm 2. The root coin represents the product of all odds contributions.

Algorithm 3: Divide-and-conquer Bernoulli factory

```

1 function FlipDCBF( $\mathcal{F}$ )
2   if  $\mathcal{F}$  contains one leaf batch then
3     return FlipLeafCoin( $\mathcal{F}$ );
4   end
5   Split  $\mathcal{F}$  into two approximately balanced subsets  $\mathcal{F}_L$  and  $\mathcal{F}_R$ ;
6    $C_L \leftarrow \text{FlipDCBF}(\mathcal{F}_L)$ ;
7    $C_R \leftarrow \text{FlipDCBF}(\mathcal{F}_R)$ ;
8   return MergeCoins( $C_L, C_R$ );
9 end

```

The role of the leaf coin is model-specific. In some applications, a leaf coin can be simulated exactly using a tractable bound. In other applications, it requires an unbiased estimator or an auxiliary construction. The exactness of the full DCBF depends on the exactness of every leaf coin and every merge operation. If a leaf coin is only approximate, then the full DCBF is also approximate.

For the normalised power prior, the DCBF can be combined with an exchange construction. The exchange step first removes the ratio $c_0(\eta)/c_0(\eta')$ by drawing θ^* from the powered historical posterior at η' . The remaining exchange odds can be factorized over historical observations or over batches. The DCBF can then be used to simulate the Barker coin associated with this factorized exchange ratio. This route is exact only when the auxiliary exchange draw is exact and the leaf coins are exact.

2.5.3.3 Computational complexity

The motivation for the DCBF is that a balanced tree can reduce the cost of simulating product-form Bernoulli probabilities. The vanilla two-coin algorithm may have exponential cost in the number of factors. By contrast, (Stumpf-Fétizon; Gonçalves, 2025) show that divide-and-conquer Bernoulli factory MCMC can have polynomial complexity under suitable regularity conditions. The precise polynomial degree depends on the construction of the leaf coins and on the availability of efficient unbiased estimators of the intractable likelihood ratio.

A simplified complexity statement is

$$\mathbb{E}[\text{Cost}] = O\left(4^\ell \rho(n, \ell)\right), \quad (2.51)$$

where ℓ is the depth of the tree and $\rho(n, \ell)$ is the expected cost of the leaf-level two-coin algorithm for a batch of size $n/2^\ell$. When the posterior concentrates at rate $n^{-1/2}$ and the proposal is optimally scaled, suitable choices of ℓ can lead to polynomial rather than

exponential complexity. This is the main theoretical advantage of the divide-and-conquer construction.

The practical message is more cautious. The DCBF is exact only if its component coins are exact. It can be computationally expensive if the merge procedure requires many repeated attempts. It also requires valid bounds or exact local constructions at the leaves. For this reason, in this thesis the DCBF is treated as a theoretically motivated exact method under bounded-support assumptions, while the exchange Metropolis–Hastings update is used as the main exact benchmark when exact powered-posterior simulation is available.

2.5.4 Pseudo-marginal Metropolis–Hastings

Pseudo-marginal Metropolis–Hastings is a modification of the Metropolis–Hastings algorithm for targets whose unnormalised density cannot be evaluated exactly, but can be estimated without bias. Let the target distribution be

$$\pi(\theta) = \frac{\gamma(\theta)}{Z}, \quad (2.52)$$

where Z is unknown and where $\gamma(\theta)$ is not available pointwise. Assume that, for each θ , we can simulate an auxiliary random variable $U \sim m_\theta(du)$ and compute a non-negative estimator $\hat{\gamma}(\theta, U)$ satisfying

$$\mathbb{E}_{U \sim m_\theta}[\hat{\gamma}(\theta, U)] = \gamma(\theta). \quad (2.53)$$

The pseudo-marginal method applies an ordinary Metropolis–Hastings algorithm on the extended space (θ, U) . The extended target is

$$\tilde{\pi}(\theta, u) = \frac{\hat{\gamma}(\theta, u)m_\theta(u)}{Z}. \quad (2.54)$$

This construction is valid because the marginal distribution of θ under $\tilde{\pi}$ is the desired target distribution. Indeed,

$$\int \tilde{\pi}(\theta, u) du = \frac{1}{Z} \int \hat{\gamma}(\theta, u)m_\theta(u) du, \quad (2.55)$$

$$= \frac{\gamma(\theta)}{Z}, \quad (2.56)$$

$$= \pi(\theta). \quad (2.57)$$

Given the current state (θ, U) , the algorithm proposes $\theta' \sim q(d\theta' | \theta)$ and independently samples $U' \sim m_{\theta'}(du')$. The proposal is accepted with probability

$$\alpha_{\text{PM}}\{(\theta, U), (\theta', U')\} = \min \left\{ 1, \frac{\hat{\gamma}(\theta', U')q(\theta | \theta')}{\hat{\gamma}(\theta, U)q(\theta' | \theta)} \right\}. \quad (2.58)$$

A crucial implementation detail is that the estimator associated with the current state must be stored and reused after rejection. If the estimator at the current state were regenerated at every iteration, the resulting chain would generally no longer be the standard pseudo-marginal chain on the extended space. Under the non-negativity and unbiasedness assumptions, the algorithm is exact in the sense that the marginal invariant distribution of θ is π . This result is the central guarantee of the pseudo-marginal approach of [Andrieu and Roberts \(2009\)](#).

The method is especially useful when $\gamma(\theta)$ contains an intractable likelihood or an intractable normalising constant ratio for which an unbiased non-negative estimator is available. In the context of normalised power priors, one may use pseudo-marginal ideas when the ratio involving $c_0(\eta)$ is replaced by a positive unbiased estimator of the corresponding ratio.

The main caveat is that unbiasedness alone is not sufficient for good computational performance. If the estimator $\hat{\gamma}(\theta, U)$ has high variance, the chain may become sticky because unusually large estimates at the current state are difficult to move away from. This phenomenon can produce high autocorrelation, poor effective sample size, and long burn-in periods even though the invariant distribution is theoretically correct.

Another caveat is that the estimator must be non-negative. Signed unbiased estimators do not directly define a probability density on the extended space and therefore do not yield the standard pseudo-marginal guarantee. Moreover, pseudo-marginal Metropolis–Hastings should not be confused with plugging an unbiased estimator into an arbitrary nonlinear acceptance probability. For instance, replacing an exact ratio R by an unbiased estimate $\hat{R}$ inside Barker’s probability $\hat{R}/(1 + \hat{R})$ is generally biased, since

$$\mathbb{E} \left[\frac{\hat{R}}{1 + \hat{R}} \right] \neq \frac{R}{1 + R}.$$

Thus, pseudo-marginal Metropolis–Hastings provides an exact route when a non-negative unbiased estimator of the unnormalised target is available, but its practical usefulness depends strongly on the stability and variance of that estimator.

2.6 Function emulation methods

A common strategy for handling an intractable function is to replace it by a statistical surrogate constructed from a finite number of evaluations at design points. In the context of the normalised power prior, the target function is $g(\eta) = \log c_0(\eta)$, whose

structural properties are established in Sections 2.4.2 and Appendix A.4. These properties motivate the use of emulators that can incorporate prior structural information rather than treating g as an arbitrary smooth function. Two families of emulators are particularly relevant here: Gaussian processes and shape-constrained additive models. Both are used in the experiments of Chapter 3, where they are assessed against the exact methods of Chapter 4.

2.6.1 Gaussian processes

A Gaussian process (GP) is a distribution over functions such that any finite collection of function values follows a multivariate normal distribution (Rasmussen; Williams, 2006). Formally, one writes $g \sim \mathcal{GP}(m, k)$, where $m : [0, 1] \rightarrow \mathbb{R}$ is the mean function and $k : [0, 1]^2 \rightarrow \mathbb{R}$ is the covariance kernel. Given evaluations $\{(\eta_i, g(\eta_i))\}_{i=1}^n$, the GP posterior is again a GP whose mean and covariance can be computed in closed form. This posterior mean serves as the surrogate $\hat{g}(\eta)$, and the posterior variance quantifies the emulation uncertainty at unobserved values of η .

A key feature of GP emulation in this setting is that, since differentiation is a linear operation, the derivative process g' is also a Gaussian process. We discuss this in more detail in Section ??.

The main practical limitation of GP emulation is its computational cost: fitting a GP to n design points requires $O(n^3)$ operations for the matrix factorisation, which becomes prohibitive as n grows. This constraint is not severe for the one-dimensional function $g : [0, 1] \rightarrow \mathbb{R}$ studied here, but would require sparse or inducing-point approximations in higher-dimensional settings.

2.6.2 Generalised additive models

Generalised additive models (GAMs) (Hastie; Tibshirani, 1986; Wood, 2017) provide a flexible semiparametric framework for regression in which the response is modelled as a sum of smooth functions of the predictors. For a scalar covariate x , a GAM takes the form

$$f(x) = \beta_0 + \sum_{j=1}^p s_j(x), \quad (2.59)$$

where each s_j is an unknown smooth function represented through a spline basis expansion. The basis functions are chosen to span a rich function space while remaining computationally tractable; common choices include thin-plate splines, cubic regression splines, and B-splines (Wood, 2017).

Model fitting proceeds by penalised likelihood maximisation, where a roughness penalty is added to the log-likelihood to control the complexity of each smooth component and prevent overfitting. The penalty is typically of the form $\lambda_j \int [s_j''(x)]^2 dx$, and the

smoothing parameter λ_j is selected by cross-validation or by restricted maximum likelihood (REML). This yields a smooth, data-adaptive approximation to f that does not require specifying a parametric form.

The main limitation of standard GAMs for function emulation is that they impose no qualitative constraints on the fitted function. In settings where the target function is known to satisfy structural properties such as monotonicity or convexity, unconstrained smoothers may produce artefacts, such as oscillations, sign violations, or locally inconsistent curvature, particularly in regions with sparse design points. This motivates the shape-constrained extensions discussed in Section [3.1.3](#).

3 Approximate inference with noisy-MCMC

3.1 Emulating the normalising constant

As established in Section 2.4, the normalising constant $c_0(\eta)$ of (2.29) cannot be evaluated analytically in general. Direct evaluation of the posterior density is therefore infeasible.

More precisely, suppose we evaluate the log-normalising constant $g(\eta) = \log c_0(\eta)$ at points $\eta_1, \dots, \eta_m \in [0, 1]$. The evaluations may be obtained analytically, via Monte Carlo estimators, thermodynamic integration, or path sampling methods. We then construct an emulator $\hat{g}(\eta)$ intended to approximate $g(\eta)$ over the entire interval. This yields the approximation $c_0(\eta) \approx \exp\{\hat{g}(\eta)\}$.

The resulting approximation can then be plugged into the posterior density, producing an approximate MCMC scheme whose computational cost is substantially reduced after the emulator is trained.

This can be interpreted as a problem of evaluating an intractable normalising constant into a nonparametric regression problem. The quality of inference then depends on the approximation properties of the emulator and on how uncertainty propagates through the posterior distribution.

3.1.1 Consequences

Although emulation methods can dramatically reduce computational cost, replacing the true normalising constant by a deterministic approximation may substantially alter the resulting posterior (Park; Haran, 2017; Park; Haran, 2020).

Suppose the true posterior density is

$$\pi(\theta, \eta \mid D, D_0) \propto \frac{L(D \mid \theta)L(D_0 \mid \theta)^\eta \pi_0(\theta)\pi_0(\eta)}{c_0(\eta)}. \quad (3.1)$$

If $c_0(\eta)$ is replaced by an approximation $\hat{c}_0(\eta)$, the resulting Markov chain targets instead

$$\tilde{\pi}(\theta, \eta \mid D, D_0) \propto \frac{L(D \mid \theta)L(D_0 \mid \theta)^\eta \pi_0(\theta)\pi_0(\eta)}{\hat{c}_0(\eta)}. \quad (3.2)$$

Consequently, approximation errors directly perturb the invariant distribution of the Markov chain. Even small systematic biases in $\hat{c}_0(\eta)$ may produce substantial

distortions in posterior summaries, credible intervals, and Bayes factors.

An important practical issue is that many emulation approaches treat $\hat{c}_0(\eta)$ as deterministic after fitting. In this case, uncertainty associated with the emulator is ignored during posterior inference. This may lead to overconfident posterior distributions and unreliable uncertainty quantification.

The problem becomes particularly severe when the approximation error varies non-uniformly across the parameter space. For instance, local interpolation errors near regions of high posterior mass may induce systematic bias in the marginal posterior distribution of η .

Moreover, unconstrained emulators may violate structural properties known to hold for the true function. For normalised power priors, the map $g(\eta) = \log c_0(\eta)$ typically satisfies smoothness and monotonicity properties inherited from the likelihood. Approximations that ignore such information may generate unrealistic artifacts, including oscillations and locally inconsistent curvature.

These issues motivate the incorporation of structural constraints into the emulation procedure.

3.1.2 Shape-constrained Gaussian process emulation

One strategy for improving the stability of the emulator is to incorporate qualitative information about the target function directly into the model.

Suppose $g(\eta) = \log c_0(\eta)$ is a Gaussian process with mean m and covariance kernel k .

Under regularity conditions, derivatives of g can be expressed in terms of moments under the power posterior distribution. For instance,

$$\frac{d}{d\eta} \log c_0(\eta) = \mathbb{E}_{\pi_\eta}[\log L(D_0 \mid \theta)]. \quad (3.3)$$

Therefore, information about monotonicity and curvature may be available analytically or estimated from Monte Carlo samples.

A flexible approach consists of modelling g as a Gaussian process, as discussed in Section 2.6.1. If

$$g \sim \mathcal{GP}(m, k), \quad (3.4)$$

then derivatives of g are also Gaussian processes because differentiation is a linear operator. In particular,

$$g' \sim \mathcal{GP}\left(\frac{d}{d\eta}m(\eta), \frac{d}{d\eta}\frac{d}{d\eta'}k(\eta, \eta')\right). \quad (3.5)$$

This allows derivative information and shape restrictions to be incorporated directly into the posterior distribution of the Gaussian process. Conditioning on derivative observations or derivative sign constraints can substantially reduce unrealistic oscillatory behaviour.

We refer to [Riihimäki and Vehtari \(2010\)](#), [Wang and Berger \(2016\)](#) and [Adame and Carvalho \(2025\)](#) for recent developments in shape-constrained emulation through Gaussian processes.

3.1.3 Shape-constrained additive models

An alternative strategy is based on shape-constrained additive models (SCAMs), which extend generalised additive models (GAMs) (as presented in Section 2.6.2) by enforcing qualitative constraints, such as monotonicity or convexity, directly on the spline basis coefficients ([Pya; Wood, 2015](#)).

In a SCAM, the target function is modelled as

$$g(\eta) = \beta_0 + \sum_{j=1}^p f_j(\eta), \quad (3.6)$$

where each f_j is represented by a spline expansion whose coefficients are constrained during estimation to satisfy the desired shape restriction.

The model is fitted by penalised likelihood under linear inequality constraints on the coefficients, using methods closely related to standard GAM estimation ([Wood, 2017](#)).

For the normalised power prior, the convexity and monotonicity of $\log c_0(\eta)$ established in Appendix A.4 provide a natural set of shape constraints that can be directly encoded in the SCAM.

By incorporating this structural information, the SCAM surrogate avoids the oscillatory artefacts that unconstrained smoothers may exhibit near regions of sparse design points. Compared with GP emulation (Section 2.6.1), SCAMs are generally faster to fit and less sensitive to kernel hyperparameter choices, at the cost of lacking a principled uncertainty quantification mechanism.

3.2 Current approaches for the NPP

The normalised power prior has received considerable attention in recent years due to its ability to incorporate historical information while preserving coherence with

the likelihood principle. However, as established in Section 2.4, the normalising constant $c_0(\eta)$ depends on η in a non-analytic manner for most models of interest, inducing a doubly intractable posterior whose evaluation requires repeated computation of $c_0(\eta)$ across MCMC iterations.

A substantial portion of the current literature therefore focuses on approximating either the normalising constant itself or ratios involving $c_0(\eta)$. The main approaches may be broadly divided into:

- (i) direct approximation of the normalising function;
- (ii) Monte Carlo estimators based on importance sampling or path sampling; and
- (iii) exact or asymptotically exact MCMC procedures.

3.2.1 Approximate inference through emulation and interpolation

One of the main practical difficulties in the NPP is that evaluating $c_0(\eta)$ repeatedly during posterior simulation may become prohibitively expensive for complex models. Consequently, several authors have proposed constructing surrogate approximations for the normalising function.

A common strategy consists of evaluating $c_0(\eta)$ or $\log c_0(\eta)$ over a finite grid of discounting parameters and then interpolating these values over the interval $\eta \in [0, 1]$. Spline-based approximations are particularly attractive in this setting because the normalising function is typically smooth as a function of η .

In the R package NPP associated with Han et al. (2023), efficient posterior computation is achieved by avoiding repeated evaluation of the normalising constant during MCMC sampling. The proposed framework exploits structural properties of the power posterior and constructs computational approximations that only require simulation from a small number of reference distributions.

More generally, spline interpolation methods approximate the log-normalising constant through basis expansions of the form

$$\hat{g}(\eta) = \sum_{k=1}^K \beta_k B_k(\eta), \quad (3.7)$$

where B_k denotes spline basis functions. The resulting approximation can then be plugged directly into the posterior distribution.

Although such methods may substantially reduce computational cost, they typically treat the approximation as deterministic after fitting. As a consequence, uncertainty associated with the approximation is not propagated into the posterior distribution.

This may lead to overconfident inference and distortions in posterior summaries when approximation errors accumulate near regions of high posterior mass.

An important limitation of the current literature is that there is still no widely adopted exact computational baseline for assessing approximation quality in the NPP setting. Most existing works compare computational approximations against other approximate procedures or against simplified models where the normalising constant is analytically available. This makes it difficult to quantify the true inferential impact of approximation errors.

The recent development of exact and quasi-exact MCMC procedures for doubly intractable distributions provides an opportunity to revisit these approximation methods under a more rigorous computational framework. In particular, exact samplers may be used as gold-standard references for evaluating the bias and stability of surrogate models.

3.2.2 Monte Carlo approaches and exact inference

Several recent works investigate Monte Carlo estimators for posterior inference under the normalised power prior. These methods typically exploit the representation of the normalising constant as an expectation under a tractable proposal distribution.

The work of [Carvalho and Ibrahim \(2021\)](#) establishes several theoretical properties of the NPP and highlights the computational challenges induced by the normalising function. In particular, the authors show that the normalised power prior is always proper whenever the initial prior is proper and prove smoothness and convexity properties of the map $\eta \mapsto \log c_0(\eta)$.

The same work also emphasizes that posterior inference under the NPP belongs to the class of doubly intractable distributions because the posterior density depends explicitly on an unknown normalising constant ratio. This observation motivates the use of advanced MCMC procedures originally developed for doubly intractable models.

More recently, [Han et al. \(2023\)](#) proposed efficient posterior sampling procedures based on importance sampling and power posterior representations. The central idea is to avoid evaluating $c_0(\eta)$ repeatedly by constructing estimators using samples generated under reference power posterior distributions.

These approaches are computationally attractive because they are relatively simple to implement and scale well to moderate-dimensional models. However, they remain approximate in general and their accuracy depends critically on the variance of the Monte Carlo estimators and on the overlap between the proposal and target distributions.

At present, the literature still lacks a systematic comparison between:

- (i) deterministic emulation procedures;

- (ii) pseudo-marginal Monte Carlo estimators; and
- (iii) exact MCMC algorithms for the NPP.

There is limited empirical evidence regarding the stationary bias induced by noisy approximations, the computational trade-offs between exact and approximate methods, and the practical benefits of incorporating structural information such as shape constraints into surrogate models. To address these questions, this work combines exact MCMC procedures based on the Exchange algorithm with approximation strategies based on functional emulation, providing a unified framework in which both exact and approximate methods can be systematically evaluated against a common computational baseline.

3.3 Function emulation for doubly intractable settings

As established in Section 2.3, doubly intractable distributions arise when the posterior density contains a normalising constant that depends on the parameter being updated and cannot be evaluated pointwise. The normalised power prior is a canonical example of this phenomenon. This section reviews the main computational strategies available in this setting, with a focus on the emulation approach of [Park and Haran \(2020\)](#) and its applicability to the NPP.

A number of MCMC methods have been proposed to address this problem, including auxiliary variable approaches such as the exchange algorithm and the double Metropolis–Hastings algorithm ([Liang, 2009](#)). While these methods target the exact posterior distribution, they require simulating auxiliary variables at each MCMC iteration, which becomes computationally expensive in complex or high-dimensional models.

An alternative strategy is proposed by [Park and Haran \(2020\)](#), who introduce a function emulation approach to approximate the intractable normalising constant. Their method replaces repeated Monte Carlo estimation of $Z(\theta)$ with a two-stage surrogate model. First, importance sampling is used to obtain noisy estimates of $Z(\theta)$ at a finite set of parameter values. Then, a Gaussian process (GP) is fitted to the log-normalising function, providing a fast interpolator that can be evaluated within an MCMC algorithm. This leads to two variants: one that emulates only the normalising function (NormEm), and another that emulates the entire likelihood (LikEm).

The main advantage of this approach is computational efficiency. By shifting the costly Monte Carlo approximations to a precomputation stage, the resulting MCMC algorithm avoids repeated simulation of auxiliary variables, yielding substantial speed-ups compared to exact or pseudo-marginal methods. Moreover, the approach is naturally parallelizable and extends beyond doubly intractable settings to more general problems involving expensive likelihood evaluations.

However, this efficiency comes at the cost of asymptotic exactness. The resulting Markov chain targets an approximate posterior distribution, with an error that depends on both the number of importance samples and the number of design points used to train the Gaussian process. While theoretical bounds on the total variation distance are provided, they do not yield explicit rates, making it difficult to calibrate the approximation in practice.

Several limitations remain. The importance sampling step can exhibit high variance in moderate dimensions, degrading the surrogate quality. The GP emulator assumes a smooth log-normalising function and ignores heteroskedasticity from Monte Carlo noise. The design points are chosen by heuristic short-chain runs, with no principled adaptive strategy. Finally, the cubic cost of GP regression limits scalability to the one-dimensional setting studied here.

These limitations suggest several directions for future research. In particular, combining surrogate-based approximations with exact correction mechanisms, such as pseudo-marginal or delayed-acceptance schemes, could potentially recover asymptotic exactness while retaining computational gains. Additionally, replacing importance sampling with more robust estimators, and incorporating adaptive or local surrogate models, may improve both stability and scalability. Such developments are particularly relevant in the context of NPP models, where efficient and reliable inference remains an open challenge.

4 Exact inference for the Normalised Power Prior

In this chapter, we develop exact Markov chain Monte Carlo methods for inference under the normalised power prior.

The full joint posterior of interest is

$$\pi(\theta, \eta \mid D, D_0) \propto L(D \mid \theta) \frac{L_0(D_0 \mid \theta)^\eta \pi_0(\theta)}{c_0(\eta)} \pi_A(\eta), \quad (4.1)$$

and the central challenge is the intractable ratio $c_0(\eta)/c_0(\eta')$ that appears in any Metropolis–Hastings update of η .

We address this challenge through three complementary constructions: the exchange algorithm (Section 4.1), the Poisson estimator (Section 4.2.1), and a pseudo-marginal approach based on importance sampling (Section 4.3).

Updating θ given (η, D, D_0) is straightforward: the full conditional

$$\pi(\theta \mid \eta, D, D_0) \propto L(D \mid \theta) L_0(D_0 \mid \theta)^\eta \pi_0(\theta) \quad (4.2)$$

does not involve $c_0(\eta)$, so standard Metropolis–Hastings (Section 2.2.2) applies directly.

For a symmetric proposal $q(\theta' \mid \theta) = q(\theta \mid \theta')$, the acceptance ratio reduces to

$$\alpha(\theta, \theta' \mid \eta) = \min \left\{ 1, \frac{L(D \mid \theta') \pi_0(\theta')}{L(D \mid \theta) \pi_0(\theta)} \left(\frac{L_0(D_0 \mid \theta')}{L_0(D_0 \mid \theta)} \right)^\eta \right\}. \quad (4.3)$$

The remainder of this chapter focuses on the harder problem of updating η .

4.1 Exchange Metropolis–Hastings

Applying the exchange algorithm of Section 2.5.1 to the NPP, the intractable quantity is $c_0(\eta)$ as defined in (2.29).

If the current value of θ is fixed and a proposal η' is generated from $q(\eta' \mid \eta)$, the exchange update draws

$$\theta^* \sim \pi_{\eta'}(\theta \mid D_0) = \frac{L_0(D_0 \mid \theta)^{\eta'} \pi_0(\theta)}{c_0(\eta')}. \quad (4.4)$$

The exchange ratio for the η update is then

$$R_{\text{ex}}(\eta, \eta'; \theta, \theta^*) = \frac{\pi_A(\eta') q(\eta \mid \eta') L_0(D_0 \mid \theta)^{\eta' - \eta}}{\pi_A(\eta) q(\eta' \mid \eta) L_0(D_0 \mid \theta^*)^{\eta' - \eta}}. \quad (4.5)$$

The ratio $c_0(\eta)/c_0(\eta')$ has disappeared. This cancellation is the main reason why the exchange algorithm is attractive for the normalised power prior.

According to the exchange algorithm framework of [Murray, Ghahramani and MacKay \(2012\)](#), the auxiliary draw from $\pi_{\eta'}(\theta \mid D_0)$ introduces a factor $1/c_0(\eta')$ into the augmented target that exactly cancels the intractable ratio (see [Appendix A.3](#) for a self-contained argument).

The exactness of the exchange algorithm depends on the ability to sample θ^* exactly from $\pi_{\eta'}(\theta \mid D_0)$. This assumption is strong. It is satisfied in conjugate Gaussian models and in some conditionally conjugate models. It is not automatically satisfied in general regression models or in non-conjugate likelihoods. If θ^* is generated by an inner MCMC chain that is not run to stationarity, the resulting algorithm is no longer exactly the exchange algorithm. In that case, the method becomes an approximate exchange or double Metropolis–Hastings algorithm.

4.2 Barker's method via Poisson estimator

Updating η is more delicate because the full conditional

$$\pi(\eta \mid \theta, D, D_0) \propto \frac{\pi_A(\eta)}{c_0(\eta)} L_0(D_0 \mid \theta)^\eta \quad (4.6)$$

involves $c_0(\eta)$ explicitly.

The Metropolis–Hastings ratio for a proposal $\eta \rightarrow \eta'$ is

$$\frac{\pi(\eta' \mid \theta)}{\pi(\eta \mid \theta)} = \frac{\pi_A(\eta')}{\pi_A(\eta)} \cdot \frac{c_0(\eta)}{c_0(\eta')} \cdot L_0(D_0 \mid \theta)^{\eta' - \eta}, \quad (4.7)$$

and the intractable ratio $c_0(\eta)/c_0(\eta')$ prevents direct evaluation.

Recalling the thermodynamic identity established in [\(2.31\)](#), integrating from η to η' gives

$$\Delta(\eta, \eta') := \log \frac{c_0(\eta')}{c_0(\eta)} = \int_{\eta}^{\eta'} \mathbb{E}_{\pi_{\tau}} [\log L_0(D_0 \mid \theta)] d\tau, \quad (4.8)$$

where $\pi_{\tau}(\theta) \propto L_0(D_0 \mid \theta)^{\tau} \pi_0(\theta)$ is the power prior at temperature τ .

The goal of the following sections is to construct a non-negative unbiased estimator of $\exp\{-\Delta(\eta, \eta')\} = c_0(\eta)/c_0(\eta')$.

4.2.1 Producing an unbiased estimator via a Poisson estimator

Retaining the notation of [\(4.8\)](#), the goal is to construct a non-negative unbiased estimator of $\exp\{-\Delta(\eta, \eta')\} = c_0(\eta)/c_0(\eta')$.

The key observation is that $\Delta(\eta, \eta')$ can be written as an expectation under a tractable proposal-specific measure. If $U \sim \text{Unif}(0, 1)$ and $\tau_{\eta, \eta'}(U) = \eta + U(\eta' - \eta)$, then

$$\Delta(\eta, \eta') = \mathbb{E}_U \left[(\eta' - \eta) \mathbb{E}_{\pi_{\tau_{\eta, \eta'}(U)}} [\log L_0(D_0 \mid \theta)] \right]. \quad (4.9)$$

This representation suggests a proposal-specific random variable. Let $\mathbf{P}_{\eta, \eta'}$ be the probability measure on $[0, 1] \times \Theta$ obtained by first drawing $U \sim \text{Unif}(0, 1)$ and then drawing $\theta \sim \pi_{\tau_{\eta, \eta'}(U)}$. Equivalently,

$$\mathbf{P}_{\eta, \eta'}(du, d\theta) = du \pi_{\tau_{\eta, \eta'}(u)}(d\theta). \quad (4.10)$$

Define

$$Z_{\eta, \eta'}(U, \theta) = (\eta' - \eta) \log L_0(D_0 \mid \theta). \quad (4.11)$$

Then

$$\mathbb{E}_{\mathbf{P}_{\eta, \eta'}}[Z_{\eta, \eta'}] = \Delta(\eta, \eta'). \quad (4.12)$$

The objective is to construct an unbiased estimator of the ratio $c_0(\eta)/c_0(\eta')$. This ratio can be written as

$$\exp\{-\Delta(\eta, \eta')\} = \frac{c_0(\eta)}{c_0(\eta')}. \quad (4.13)$$

Assume that there exist deterministic constants $a_{\eta, \eta'}$ and $b_{\eta, \eta'}$ such that

$$a_{\eta, \eta'} \leq Z_{\eta, \eta'}(U, \theta) \leq b_{\eta, \eta'} \quad \mathbf{P}_{\eta, \eta'}\text{-almost surely}. \quad (4.14)$$

Set

$$c_{\eta, \eta'} = b_{\eta, \eta'}, \quad \lambda_{\eta, \eta'} = b_{\eta, \eta'} - a_{\eta, \eta'}. \quad (4.15)$$

Draw $K \sim \text{Poisson}(\lambda_{\eta, \eta'})$. Conditionally on K , draw

$$(U_j, \theta_j)_{j=1}^K \stackrel{\text{iid}}{\sim} \mathbf{P}_{\eta, \eta'}. \quad (4.16)$$

The Poisson estimator is

$$\hat{I}(\eta, \eta') = \exp\{\lambda_{\eta, \eta'} - c_{\eta, \eta'}\} \prod_{j=1}^K \frac{c_{\eta, \eta'} - Z_{\eta, \eta'}(U_j, \theta_j)}{\lambda_{\eta, \eta'}}. \quad (4.17)$$

By construction, each factor satisfies

$$0 \leq \frac{c_{\eta, \eta'} - Z_{\eta, \eta'}(U_j, \theta_j)}{\lambda_{\eta, \eta'}} \leq 1. \quad (4.18)$$

Therefore $\hat{I}(\eta, \eta')$ is non-negative. Moreover, $\hat{I}(\eta, \eta')$ is an unbiased estimator of $c_0(\eta)/c_0(\eta')$; the proof is given in Appendix A.2.

The efficiency of the estimator depends strongly on the envelope width $\lambda_{\eta, \eta'} = b_{\eta, \eta'} - a_{\eta, \eta'}$. Large envelopes lead to large Poisson counts and high estimator variability. Therefore, the practical usefulness of the construction depends on obtaining deterministic but sufficiently tight proposal-specific bounds for $Z_{\eta, \eta'}$.

It is important to distinguish this unbiased estimator from an exact Barker accept–reject mechanism. Although $\hat{I}(\eta, \eta')$ is unbiased for $c_0(\eta)/c_0(\eta')$, substituting a noisy estimate of the likelihood ratio directly into the Barker probability does not generally produce an exact Barker transition. Exact Barker implementations require simulating the accept–reject event itself with the correct marginal probability, typically through a Bernoulli-factory or two-coin construction.

4.2.1.1 Practical implementation

The integral $\Delta(\eta, \eta')$ in (4.8) is estimated by Monte Carlo.

Drawing $\tau_k \stackrel{\text{iid}}{\sim} \text{Unif}(\eta, \eta')$ and, conditionally, $\tilde{\theta}_{m,k} \sim \pi_{\tau_k}$ for $m = 1, \dots, M$ and $k = 1, \dots, K$, gives the double Monte Carlo estimator

$$\hat{\Delta}(\eta, \eta') = \frac{(\eta' - \eta)}{MK} \sum_{k=1}^K \sum_{m=1}^M \log L_0(D_0 \mid \tilde{\theta}_{m,k}). \quad (4.19)$$

This is an unbiased estimator of $\Delta(\eta, \eta')$ by the law of total expectation.

In the single-draw case ($K = M = 1$), this reduces to

$$\hat{\Delta}(\eta, \eta') = (\eta' - \eta) \log L_0(D_0 \mid \tilde{\theta}_*), \quad \tau_* \sim \text{Unif}(\eta, \eta'), \quad \tilde{\theta}_* \sim \pi_{\tau_*}, \quad (4.20)$$

which can be used in the Poisson estimator construction of Section 4.2.1.

Note that only the mean of $Z_{\eta, \eta'}$, not its variance, enters the unbiasedness argument; the variance governs the efficiency of the estimator through the expected Poisson count $\lambda_{\eta, \eta'}$.

4.2.2 Local envelopes for the Poisson estimator

The preceding construction does not require a single global envelope that is valid uniformly over all possible proposals. Instead, for each proposed move $\eta \rightarrow \eta'$, it is enough to obtain deterministic bounds for the proposal-specific random variable

$$Z_{\eta, \eta'}(U, \theta) = (\eta' - \eta) \log L_0(D_0 \mid \theta), \quad (U, \theta) \sim \mathbf{P}_{\eta, \eta'}. \quad (4.21)$$

Suppose that the historical log-likelihood satisfies

$$\ell_{\min} \leq \log L_0(D_0 \mid \theta) \leq \ell_{\max} \quad \text{for all } \theta \in \Theta. \quad (4.22)$$

Then, for a fixed proposal $\eta \rightarrow \eta'$, define

$$a_{\eta, \eta'} = \min\{(\eta' - \eta)\ell_{\min}, (\eta' - \eta)\ell_{\max}\}. \quad (4.23)$$

Similarly, define

$$b_{\eta, \eta'} = \max\{(\eta' - \eta)\ell_{\min}, (\eta' - \eta)\ell_{\max}\}. \quad (4.24)$$

The resulting envelope width is

$$b_{\eta, \eta'} - a_{\eta, \eta'} = |\eta' - \eta|(\ell_{\max} - \ell_{\min}). \quad (4.25)$$

Hence a natural local choice is

$$c_{\eta, \eta'} = b_{\eta, \eta'}, \quad \lambda_{\eta, \eta'} = |\eta' - \eta|(\ell_{\max} - \ell_{\min}). \quad (4.26)$$

This choice makes the expected number of Poisson factors proportional to the proposal distance $|\eta' - \eta|$. Consequently, local random-walk proposals in η lead to smaller expected computational cost than large independent proposals.

For unbounded parameter spaces, such as Gaussian models with $\Theta = \mathbb{R}^d$, the lower bound $\ell_{\min}$ is typically $-\infty$. In that case, the exact Poisson construction is not directly applicable without additional structure. One route is to truncate the prior to a compact set $\Theta_B \subset \Theta$, replacing π_0 with

$$\pi_{0,B}(\theta) \propto \pi_0(\theta) \mathbf{1}_{\Theta_B}(\theta). \quad (4.27)$$

For the Gaussian historical likelihood with known covariance, the log-likelihood is a concave quadratic in θ , attaining its maximum at the historical MLE and its minimum over any compact box at the boundary. These properties yield deterministic bounds $[\ell_{\min}, \ell_{\max}]$ on Θ_B and thus a valid local envelope.

Remark (Scope of the Gaussian experiment). The Gaussian experiment of Section 5.2.1 uses a Gaussian prior with full support $\Theta = \mathbb{R}^3$. The Poisson estimator is therefore not applied in that experiment: the methods evaluated there are Exchange MH, Prior IS, MLE-centred IS, and Path Sampling PMMH. The local envelope construction above is provided as theoretical groundwork for applications where the prior is compactly supported or can be practically truncated without loss of inferential fidelity.

4.2.3 Divide-and-conquer factorization

The computational difficulty of the Poisson estimator increases with the variability of the historical log-likelihood. Since the historical likelihood factorizes over observations, we can write

$$\log L_0(D_0 | \theta) = \sum_{i=1}^{N_0} \log L_i(d_{0i} | \theta). \quad (4.28)$$

The width of a global envelope for $Z_{\eta, \eta'} = (\eta' - \eta) \log L_0(D_0 | \theta)$ typically grows with N_0 . This motivates a divide-and-conquer construction.

Let $D_{0,1}, \dots, D_{0,G}$ be a partition of the historical data. Let $\mathcal{I}_g$ denote the set of observation indices in block g . Define the block likelihood by

$$L_{0,g}(D_{0,g} | \theta) = \prod_{i \in \mathcal{I}_g} L_i(d_{0i} | \theta). \quad (4.29)$$

Then the full historical likelihood satisfies

$$L_0(D_0 \mid \theta) = \prod_{g=1}^G L_{0,g}(D_{0,g} \mid \theta). \quad (4.30)$$

Equivalently, the full historical log-likelihood satisfies

$$\log L_0(D_0 \mid \theta) = \sum_{g=1}^G \log L_{0,g}(D_{0,g} \mid \theta). \quad (4.31)$$

For each block g , define the blockwise random variable

$$Z_{g,\eta,\eta'}(U, \theta) = (\eta' - \eta) \log L_{0,g}(D_{0,g} \mid \theta). \quad (4.32)$$

Assume that deterministic bounds $a_{g,\eta,\eta'}$ and $b_{g,\eta,\eta'}$ are available such that

$$a_{g,\eta,\eta'} \leq Z_{g,\eta,\eta'}(U, \theta) \leq b_{g,\eta,\eta'} \quad \mathbf{P}_{\eta,\eta'}\text{-almost surely.} \quad (4.33)$$

Set

$$c_{g,\eta,\eta'} = b_{g,\eta,\eta'}, \quad \lambda_{g,\eta,\eta'} = b_{g,\eta,\eta'} - a_{g,\eta,\eta'}. \quad (4.34)$$

A blockwise Poisson estimator can then be constructed as

$$\hat{I}_g(\eta, \eta') = \exp\{\lambda_{g,\eta,\eta'} - c_{g,\eta,\eta'}\} \prod_{j=1}^{K_g} \frac{c_{g,\eta,\eta'} - Z_{g,\eta,\eta'}(U_{g,j}, \theta_{g,j})}{\lambda_{g,\eta,\eta'}}, \quad (4.35)$$

where $K_g \sim \text{Poisson}(\lambda_{g,\eta,\eta'})$. The variables $(U_{g,j}, \theta_{g,j})$ are sampled independently from the same proposal-specific measure $\mathbf{P}_{\eta,\eta'}$. The blockwise estimator satisfies

$$\mathbb{E}[\hat{I}_g(\eta, \eta')] = \exp\{-\Delta_g(\eta, \eta')\}, \quad (4.36)$$

where

$$\Delta_g(\eta, \eta') = \mathbb{E}_{\mathbf{P}_{\eta,\eta'}} [Z_{g,\eta,\eta'}(U, \theta)]. \quad (4.37)$$

Proof. If the block estimators are independent, the divide-and-conquer product estimator is

$$\hat{I}_{\text{dc}}(\eta, \eta') = \prod_{g=1}^G \hat{I}_g(\eta, \eta'). \quad (4.38)$$

By independence, its expectation is

$$\mathbb{E}[\hat{I}_{\text{dc}}(\eta, \eta')] = \prod_{g=1}^G \mathbb{E}[\hat{I}_g(\eta, \eta')] \quad (4.39)$$

$$= \prod_{g=1}^G \exp\{-\Delta_g(\eta, \eta')\} \quad (4.40)$$

$$= \exp\left\{-\sum_{g=1}^G \Delta_g(\eta, \eta')\right\}. \quad (4.41)$$

Since

$$\sum_{g=1}^G \Delta_g(\eta, \eta') = \Delta(\eta, \eta'), \quad (4.42)$$

we obtain

$$\mathbb{E}[\hat{I}_{\text{dc}}(\eta, \eta')] = \exp\{-\Delta(\eta, \eta')\}. \quad (4.43)$$

Thus, the divide-and-conquer product estimator is unbiased for $c_0(\eta)/c_0(\eta')$. $\square$

The main advantage of this construction is variance and cost control. Instead of using one global envelope with width $\lambda_{\eta, \eta'}$, the method uses several smaller envelopes with widths $\lambda_{g, \eta, \eta'}$. The expected number of Poisson factors is $\sum_{g=1}^G \lambda_{g, \eta, \eta'}$. When the blockwise envelopes are much tighter than the full-data envelope, the product construction can be substantially more stable than the global estimator.

This product estimator should not be confused with a genuine recursive divide-and-conquer Bernoulli factory. The estimator $\hat{I}_{\text{dc}}(\eta, \eta')$ is an unbiased estimator of a normalising-constant ratio. However, using it as a noisy plug-in inside the Barker probability does not generally yield an exact Barker transition. A true DCBF Barker transition recursively constructs the accept–reject coin itself. At the leaves, one simulates smaller Bernoulli coins associated with subsets of likelihood factors. At internal nodes, these coins are merged recursively until a single coin with the Barker acceptance probability is obtained at the root. Therefore, the recursive DCBF is a coin construction, not merely a product of random ratio estimates.

4.3 Pseudo-marginal via importance sampling

An alternative to constructing an unbiased estimator of the ratio $c_0(\eta)/c_0(\eta')$ is to directly estimate the normalising constant $c_0(\eta)$ itself and form the ratio from two independent estimates.

The simplest such approach is the *prior Monte Carlo estimator*, which exploits the representation

$$c_0(\eta) = \mathbb{E}_{\theta \sim \pi_0}[L_0(\theta \mid D_0)^\eta].$$

Samples drawn directly from the prior π_0 are cheap to obtain and require no gradient information or auxiliary computations. The main limitation is that the prior and the power posterior $\pi_\eta(\theta) \propto L_0(\theta \mid D_0)^\eta \pi_0(\theta)$ can be poorly aligned, particularly when the historical data are informative, leading to high variance in the importance weights and an inefficient estimator.

A more principled alternative is *importance sampling* centered at the maximum likelihood estimate. Writing the normalising constant as

$$c_0(\eta) = \int_{\theta \in \Theta} L_0(\theta \mid D_0)^\eta \pi_0(\theta) \, \mathrm{d}\theta = \mathbb{E}_{\theta \sim \pi_{\text{MLE}}} \left[L_0(\theta \mid D_0)^\eta \frac{\pi_0(\theta)}{\pi_{\text{MLE}}(\theta)} \right],$$

one draws samples from a proposal π_{MLE} centered at the MLE and reweights by the ratio $\pi_0(\theta)/\pi_{\text{MLE}}(\theta)$. The MLE-centred proposal concentrates mass where the likelihood is large, reducing variance relative to the prior estimator when the historical data are informative. The cost is the need to compute or approximate the MLE, which may itself require numerical optimisation.

5 Experiments

5.1 Assessing exact and approximate samplers

The primary objective of this chapter is to compare exact and approximate inference procedures for the normalised power prior under controlled settings where the normalising constant is analytically available.

This allows us to evaluate not only computational efficiency, but also the inferential distortion introduced by noisy approximations of the normalising constant ratio.

A central distinction throughout the experiments is between:

- (i) exact MCMC algorithms targeting the true posterior distribution; and
- (ii) pseudo-marginal or approximation-based methods whose stationary distribution depends on Monte Carlo estimators or surrogate approximations.

Rather than comparing methods exclusively through asymptotic guarantees, we focus on their finite-sample practical behaviour. In particular, we evaluate posterior agreement with the analytic baseline, computational efficiency, acceptance rates, Monte Carlo variability, and stability of the normalising-constant estimators.

Since the posterior distribution of η is one-dimensional, we use the empirical Wasserstein distance relative to the exact analytic sampler as the primary discrepancy measure. Unlike density-based divergences, the Wasserstein distance does not require kernel density estimation and remains numerically stable even when posterior samples concentrate near the boundary of the parameter space.

5.1.1 Using exact methods as baseline

The Exchange MH algorithm described in Section 4.1 serves as the primary exact reference throughout the experiments. Where the power prior $\pi_{\eta'}(\theta \mid D_0)$ is available in closed form (Gaussian and Normal–Inverse–Gamma models), the exchange update is exact and computationally efficient.

The importance sampling estimators δ_a (prior IS) and δ_b (MLE-centred IS), introduced in Section 4.3, are used as pseudo-marginal baselines. They provide non-negative unbiased estimates of $c_0(\eta)$ and can be embedded within a Metropolis–Hastings scheme, yielding exact pseudo-marginal chains under the conditions of [Andrieu and Roberts \(2009\)](#).

5.2 Convergence to the baseline

5.2.1 Normal distribution with known covariance

We evaluate the proposed samplers under a multivariate Gaussian model with known covariance structure. The parameter dimension is fixed at $p = 3$, with

$$\Sigma = \mathbf{I}_3, \quad \Sigma_0 = 2\mathbf{I}_3. \quad (5.1)$$

A Gaussian prior is assigned to θ ,

$$\theta \sim \mathcal{N}_3(\mathbf{0}, \Sigma_0), \quad (5.2)$$

and the power parameter receives a uniform prior,

$$\eta \sim \text{Beta}(1, 1). \quad (5.3)$$

Historical and current datasets each contain $N_0 = N = 50$ observations.

Three discrepancy scenarios are considered. In the small-discrepancy regime, the historical and current datasets are generated from nearby means. The moderate-discrepancy regime introduces a moderate shift between the datasets. The large-discrepancy regime generates substantial disagreement between the historical and current information. The purpose of this experiment is not only to compare posterior distributions, but also to assess the computational budget required for each method to approximate the exact posterior distribution. To accomplish this, we use a convergence-to-baseline protocol.

The analytic Metropolis–Hastings sampler is used as the reference method. Since the Gaussian normalised power prior admits a closed-form normalising constant ratio, this sampler directly evaluates the exact posterior density. Longer chains are used for the analytic reference in order to reduce Monte Carlo variability and provide a stable empirical baseline for posterior comparisons. The remaining methods are evaluated according to how efficiently they approach the analytic reference posterior. We compare four methods: Exchange MH, Prior IS, MLE-centred IS, and Path Sampling PMMH.

For each method, chains are accumulated until the one-dimensional Wasserstein distance between the empirical posterior distribution of η and the analytic reference posterior falls below a fixed tolerance threshold. We additionally monitor the Wasserstein distance for $\|\theta\|_2$ in order to assess agreement in the parameter space. The primary performance metric is therefore the computational effort required to achieve a prescribed approximation accuracy. We report the Wasserstein distance to the analytic baseline, total runtime, effective sample size, posterior means and standard deviations, and acceptance rates.

A central distinction between the methods is how the ratio of normalising constants is handled. The Exchange MH algorithm avoids evaluating the normalising constant

entirely by introducing an auxiliary draw from the power prior distribution at the proposed value of η . In the present Gaussian experiment, sampling from the power prior can be performed exactly in closed form, making the exchange update computationally efficient while preserving exactness.

The Prior IS and MLE-centred IS methods instead estimate the normalising constant ratio through importance sampling constructions. Both methods remain exact pseudo-marginal algorithms when embedded within the extended-state Metropolis–Hastings framework, provided the estimator is nonnegative and unbiased. The difference between the two approaches lies primarily in variance reduction. The MLE-centred proposal attempts to concentrate samples in regions of high likelihood in order to stabilize the importance weights.

The Path Sampling PMMH method approximates the log normalising constant ratio through thermodynamic integration along a temperature path. Although effective in some settings, the resulting estimator is not guaranteed to satisfy the conditions required for exact pseudo-marginal inference and is therefore treated as an approximate baseline.




Figure 1 – Posterior density of η across discrepancy regimes for the Gaussian normalised power prior experiment.




Figure 2 – Posterior density of $\|\theta\|_2$ across discrepancy regimes for the Gaussian normalised power prior experiment.

Scenario	Method	Batches	Time (s)	W_η	$W_{\ \theta\ }$	ESS_η	$\text{ESS}_{\ \theta\ }$	Acc. η
Small	Analytic MH (exact reference)	5	103.5	0.000	0.000	552	2327	0.978
	Exchange MH	1	9.4	0.005	0.006	63	226	0.953
	Prior IS	10	2271.1	0.072	0.002	517	2111	0.965
	MLE-centered IS	7	4515.5	0.005	0.003	387	1421	0.978
	Path-sampling PMMH	6	5717.8	0.008	0.003	333	1234	0.978
Moderate	Analytic MH (exact reference)	5	94.4	0.000	0.000	1046	2677	0.907
	Exchange MH	1	9.4	0.008	0.004	123	266	0.847
	Prior IS	1	209.7	0.006	0.007	175	362	0.909
	MLE-centered IS	4	2424.3	0.009	0.003	229	945	0.907
	Path-sampling PMMH	1	922.8	0.008	0.011	161	272	0.911
Large	Analytic MH (exact reference)	5	90.7	0.000	0.000	18361	3474	0.749
	Exchange MH	1	8.9	0.001	0.007	1474	401	0.448
	Prior IS	1	168.8	0.001	0.006	872	390	0.748
	MLE-centered IS	1	484.9	0.000	0.004	1923	350	0.748
	Path-sampling PMMH	1	739.1	0.000	0.004	1954	398	0.753

Table 1 – Experiment 1 convergence diagnostics. Batches: number of batches of 4 chains $\times$ 5000 iterations required to reach $W_1(\eta) < 0.01$. Wasserstein distances computed against the analytic MH reference.

5.2.1.1 Discussion of results

The results reveal a clear distinction between methods that directly target the exact posterior distribution and methods that rely on noisier approximations of the normalising constant ratio.

First, the Exchange MH algorithm is nearly indistinguishable from the analytic reference posterior across all discrepancy regimes. This confirms empirically that the exchange construction correctly removes the intractable normalising constant ratio while preserving the target posterior distribution. Moreover, because auxiliary draws from the power prior are available analytically in the Gaussian setting, the exchange update remains computationally efficient despite being exact.

Scenario	Method	$\bar{\eta}$	$\text{sd}(\eta)$	$\text{MCSE}(\eta)$	$\ \theta\ $	$\text{sd}(\ \theta\)$	$\text{MCSE}(\ \theta\)$
Small	Analytic MH (exact reference)	0.623	0.241	0.0103	0.223	0.089	0.0018
	Exchange MH	0.621	0.246	0.0311	0.228	0.087	0.0058
	Prior IS	0.695	0.238	0.0105	0.222	0.089	0.0019
	MLE-centered IS	0.626	0.244	0.0124	0.225	0.090	0.0024
	Path-sampling PMMH	0.615	0.243	0.0133	0.226	0.088	0.0025
Moderate	Analytic MH (exact reference)	0.174	0.138	0.0043	0.291	0.115	0.0022
	Exchange MH	0.182	0.147	0.0132	0.291	0.113	0.0069
	Prior IS	0.173	0.128	0.0097	0.287	0.110	0.0058
	MLE-centered IS	0.183	0.160	0.0106	0.291	0.118	0.0038
	Path-sampling PMMH	0.169	0.124	0.0098	0.296	0.105	0.0064
Large	Analytic MH (exact reference)	0.031	0.024	0.0002	0.268	0.107	0.0018
	Exchange MH	0.032	0.025	0.0006	0.264	0.103	0.0051
	Prior IS	0.032	0.026	0.0009	0.262	0.105	0.0053
	MLE-centered IS	0.032	0.024	0.0005	0.270	0.109	0.0058
	Path-sampling PMMH	0.031	0.024	0.0005	0.265	0.108	0.0054

Table 2 – Experiment 1 posterior summaries (Normal-Normal NPP).

The Prior IS and MLE-centred IS methods both exhibit sensitivity to the discrepancy regime. Under low discrepancy, where the posterior distribution of η is diffuse, the variance of the importance weights becomes more pronounced and larger computational effort is required to achieve the target Wasserstein tolerance. The MLE-centred proposal generally improves stability relative to Prior IS by concentrating importance samples near regions of higher likelihood support.

As the discrepancy increases, the posterior distribution of η concentrates near zero. In this regime, all methods become substantially easier to stabilize because the relevant range of powered likelihoods becomes smaller. Consequently, the Wasserstein distances decrease rapidly and the differences between methods become less pronounced.

The Path Sampling PMMH method displays the largest variability across discrepancy scenarios. Because the thermodynamic integration estimator depends on Monte Carlo approximations along an entire temperature path, estimator variability accumulates and may degrade mixing efficiency. This effect is especially visible in diffuse posterior regimes.

Overall, the experiments demonstrate that exact inference for the normalised power prior is computationally feasible through the Exchange algorithm, even in settings where direct evaluation of the normalising constant ratio would otherwise be computationally prohibitive. The exchange sampler additionally provides a principled exact baseline against which approximate pseudo-marginal and future surrogate-based methods may be assessed.

5.2.2 Linear regression with Normal-Inverse-Gamma priors

We next consider a Gaussian linear regression model with unknown variance. This experiment is designed to evaluate the exchange update in a setting where the normalised

power prior remains analytically tractable, but the normalising constant is no longer a simple linear function of the power parameter. Let $D_0 = (\mathbf{X}_0, \mathbf{y}_0)$ denote the historical data and let $D = (\mathbf{X}, \mathbf{y})$ denote the current data. We assume

$$\mathbf{y}_0 = \mathbf{X}_0 \boldsymbol{\beta} + \boldsymbol{\varepsilon}_0, \quad \mathbf{y} = \mathbf{X} \boldsymbol{\beta} + \boldsymbol{\varepsilon}, \quad (5.4)$$

where $\boldsymbol{\varepsilon}_0 \sim \mathcal{N}(\mathbf{0}, \sigma^2 \mathbf{I}_{n_0})$ and $\boldsymbol{\varepsilon} \sim \mathcal{N}(\mathbf{0}, \sigma^2 \mathbf{I}_n)$. The variance σ^2 is unknown. The initial prior is Normal-Inverse-Gamma,

$$\boldsymbol{\beta} \mid \sigma^2 \sim \mathcal{N}(\boldsymbol{\mu}_0, \sigma^2 \mathbf{V}_0), \quad \sigma^2 \sim \text{Inv-Gamma}(a, b). \quad (5.5)$$

For a fixed value of η , the power prior is given by

$$\pi_\eta(\boldsymbol{\beta}, \sigma^2 \mid D_0) = \frac{L(D_0 \mid \boldsymbol{\beta}, \sigma^2)^\eta \pi_0(\boldsymbol{\beta}, \sigma^2)}{c_0(\eta)}. \quad (5.6)$$

The normalising constant is

$$c_0(\eta) = \int_0^\infty \int_{\mathbb{R}^p} L(D_0 \mid \mathbf{u}, v)^\eta \pi_0(\mathbf{u}, v) d\mathbf{u} dv. \quad (5.7)$$

In this model, $c_0(\eta)$ is available in closed form by conjugacy. This allows us to construct an analytic reference chain for η and to validate algorithms that do not directly evaluate $c_0(\eta)$. The marginal posterior density of η is proportional to

$$\pi(\eta \mid D_0, D) \propto \frac{\pi_A(\eta) \Gamma(n/2 + \nu_0) |\Lambda_0|^{1/2} H_0(\eta)^{\nu_0}}{|\Lambda|^{1/2} \Gamma(\nu_0) H(\eta)^{\nu_0 + n/2}}. \quad (5.8)$$

The quantities ν_0 , Λ_0 , $H_0(\eta)$, Λ , and $H(\eta)$ are defined in Appendix B.2. Given the current state $(\boldsymbol{\beta}, \sigma^2, \eta)$, we propose η' and draw an auxiliary value

$$(\boldsymbol{\beta}^*, \sigma^{2*}) \sim \pi_{\eta'}(\boldsymbol{\beta}, \sigma^2 \mid D_0). \quad (5.9)$$

The exchange log-ratio for the η update is

$$\begin{aligned} \log R_{\text{ex}} &= (\eta' - \eta) \log L(D_0 \mid \boldsymbol{\beta}, \sigma^2) + (\eta - \eta') \log L(D_0 \mid \boldsymbol{\beta}^*, \sigma^{2*}) \\ &\quad + \log \frac{\pi_A(\eta')}{\pi_A(\eta)} + \log \frac{q(\eta \mid \eta')}{q(\eta' \mid \eta)}. \end{aligned} \quad (5.10)$$

The proposal is then accepted with probability $\min\{1, \exp(\log R_{\text{ex}})\}$. When the auxiliary draw from $\pi_{\eta'}(\boldsymbol{\beta}, \sigma^2 \mid D_0)$ is exact, this update has the correct invariant distribution for the normalised power prior. This follows from the exchange algorithm for doubly-intractable distributions, where an auxiliary draw from the proposed normalised model cancels the unknown ratio of normalising constants in the Metropolis–Hastings acceptance probability (Murray; Ghahramani; MacKay, 2012). The exchange algorithm therefore provides an exact update for η under the powered-prior simulation assumption. This assumption is satisfied in the present experiment because the power prior is Normal-Inverse-Gamma.

We compare three methods. The first method is an analytic Metropolis–Hastings reference that evaluates Equation (5.8). The second method is the exchange Metropolis–Hastings update in Equation (5.10). The third method is a path-sampling pseudo-marginal method based on a Monte Carlo approximation of the normalising constant ratio. The path-sampling method is included as a computational baseline and is classified according to the exactness of the estimator used in its acceptance ratio. If the estimator is not non-negative and unbiased for the required marginal target, the method is treated as approximate rather than exact. This distinction follows the pseudo-marginal principle, where exactness requires a non-negative unbiased estimator of the target density on an extended space (Andrieu; Roberts, 2009).

We generate three discrepancy scenarios between the historical and current regression coefficients. In the small-discrepancy scenario, the current coefficient vector is close to the historical coefficient vector. In the moderate-discrepancy scenario, the current coefficient vector is moderately shifted. In the large-discrepancy scenario, the current coefficient vector is substantially shifted. The expected qualitative behavior is that the posterior distribution of η moves toward smaller values as the discrepancy increases. This reflects reduced borrowing from the historical data when the historical and current data-generating mechanisms are less compatible.

The main diagnostic is the one-dimensional Wasserstein distance from each method to the analytic reference. We report this distance separately for η , $\|\beta\|_2$, and σ^2 . We also report posterior means, posterior standard deviations, effective sample sizes, acceptance rates, and runtime. The norm $\|\beta\|_2$ summarizes the magnitude of the regression coefficient vector, while σ^2 is reported separately because it represents the residual variance and should not be combined with β into a single geometric average.




Figure 3 – Posterior density of η for the linear regression NPP experiment across discrepancy scenarios.




Figure 4 – Posterior density of $\|\beta\|_2$ for the linear regression NPP experiment across discrepancy scenarios.




Figure 5 – Posterior density of σ^2 for the linear regression NPP experiment across discrepancy scenarios.

Scenario	Method	W_η	$W_{\ \beta\ }$	W_σ	ESS_η	Acc. η	Time (s)
Small	Analytic MH	0.000	0.000	0.000	195	0.928	3.1
Small	Exchange MH	0.019	0.001	0.002	222	0.902	1.6
Small	Path PMMH	0.022	0.001	0.001	199	0.913	31.6
Moderate	Analytic MH	0.000	0.000	0.000	717	0.859	3.2
Moderate	Exchange MH	0.003	0.003	0.002	557	0.797	1.6
Moderate	Path PMMH	0.002	0.001	0.002	797	0.837	32.5
Large	Analytic MH	0.000	0.000	0.000	682	0.863	3.1
Large	Exchange MH	0.001	0.003	0.009	577	0.824	1.6
Large	Path PMMH	0.000	0.003	0.003	694	0.858	32.1

Table 3 – Diagnostic comparison for the linear regression NPP experiment with Normal-Inverse-Gamma priors. Wasserstein distances are computed against the analytic MH baseline within each discrepancy scenario.

Scenario	Method	$\bar{\eta}$	$\text{sd}(\eta)$	$\text{MCSE}(\eta)$	$\ \beta\ $	$\text{sd}(\ \beta\)$	$\text{MCSE}(\ \beta\)$	$\bar{\sigma}$	$\text{sd}(\sigma)$	$\text{MCSE}(\sigma)$
Small	Analytic MH	0.697	0.208	0.0149	1.227	0.076	0.0006	0.820	0.101	0.0008
Small	Exchange MH	0.679	0.217	0.0146	1.228	0.075	0.0006	0.818	0.101	0.0008
Small	Path PMMH	0.705	0.227	0.0161	1.229	0.075	0.0006	0.821	0.101	0.0008
Moderate	Analytic MH	0.077	0.049	0.0018	1.984	0.102	0.0014	0.862	0.137	0.0015
Moderate	Exchange MH	0.080	0.051	0.0021	1.981	0.101	0.0011	0.863	0.137	0.0014
Moderate	Path PMMH	0.078	0.047	0.0017	1.983	0.102	0.0012	0.863	0.138	0.0015
Large	Analytic MH	0.007	0.005	0.0002	3.912	0.102	0.0009	0.863	0.142	0.0012
Large	Exchange MH	0.007	0.005	0.0002	3.915	0.101	0.0008	0.854	0.140	0.0014
Large	Path PMMH	0.007	0.005	0.0002	3.915	0.103	0.0008	0.860	0.141	0.0013

Table 4 – Posterior summaries for the NIG linear regression NPP experiment.

5.2.2.1 Discussion of results

The most striking finding of this experiment is the close agreement between the Exchange Metropolis–Hastings sampler and the Path Sampling PMMH method across all three discrepancy scenarios. Both methods produce posterior distributions of η , $\|\beta\|_2$, and σ^2 that are nearly indistinguishable from the analytic reference, as measured by the Wasserstein distance. This is a non-trivial result, because the two methods operate through fundamentally different mechanisms.

The exchange update avoids evaluating $c_0(\eta)$ entirely by introducing an exact auxiliary draw from the powered Normal–Inverse–Gamma prior. The path-sampling method instead approximates the log-ratio $\log c_0(\eta')/c_0(\eta)$ through Monte Carlo integration along a temperature path, and is therefore only asymptotically exact when the number of integration points is large. The empirical equivalence of the two methods in this setting suggests that, for the Normal–Inverse–Gamma model, the path-sampling estimator achieves sufficient accuracy at the number of integration points used, and its residual bias is too small to produce discernible differences in posterior summaries.

This observation is practically important: it implies that the path-sampling pseudo-marginal approach can serve as a reliable and computationally competitive alternative to the exchange algorithm in settings where the power prior is conjugate and exact simulation is inexpensive. At the same time, the analytic reference confirms that both methods are correctly targeting the true posterior, lending further empirical credibility to the exchange correctness argument of Appendix A.3.

The behaviour across discrepancy regimes follows the same qualitative pattern as in the Gaussian experiment of Section 5.2.1. As the discrepancy between historical and current data increases, the posterior of η shifts towards smaller values, reflecting the adaptive borrowing property of the normalised power prior. All three methods capture this behaviour faithfully, with no systematic distortion of the marginal posterior of σ^2 or β across regimes.

5.2.3 Logistic regression with surrogate normalising constants

This experiment considers a logistic regression model for which the normalising constant of the normalised power prior is not available in closed form. The objective is to evaluate whether surrogate models can accurately emulate the function $c_0(\eta)$ and thereby provide a computationally efficient alternative to exact inference.

Let $D_0 = (\mathbf{X}_0, \mathbf{y}_0)$ denote the historical dataset and $D = (\mathbf{X}, \mathbf{y})$ denote the current dataset. We assume

$$y_{0i} \mid \boldsymbol{\beta} \sim \text{Bernoulli}(p_{0i}), \quad \text{logit}(p_{0i}) = \mathbf{x}_{0i}^\top \boldsymbol{\beta}, \quad (5.11)$$

and

$$y_i \mid \boldsymbol{\beta} \sim \text{Bernoulli}(p_i), \quad \text{logit}(p_i) = \mathbf{x}_i^\top \boldsymbol{\beta}. \quad (5.12)$$

The regression coefficient vector has prior distribution $\boldsymbol{\beta} \sim \mathcal{N}(\mathbf{0}, \mathbf{I}_p)$, with $p = 3$.

Unlike the Gaussian and Normal–Inverse–Gamma examples considered previously, no closed-form expression is available for $c_0(\eta)$. Consequently, direct Metropolis–Hastings updates involving the ratio $c_0(\eta')/c_0(\eta)$ are not feasible.

To obtain an exact reference posterior, we employ the Exchange algorithm. For a proposed value η' , an auxiliary draw $\boldsymbol{\beta}^* \sim \pi_{\eta'}(\boldsymbol{\beta} \mid D_0)$ is generated using Stan ([Carpenter et al., 2017](#)). The resulting exchange acceptance ratio depends only on likelihood evaluations and does not require explicit computation of the normalising constant ratio. Under exact simulation from the power prior, the resulting Markov chain targets the correct normalised power prior posterior distribution. The remaining methods replace $c_0(\eta)$ by an emulator constructed from a finite set of training evaluations.

We consider three surrogate models. The first is a spline emulator based on smoothing splines fitted to estimates of $\log c_0(\eta)$ over a grid of power parameters. The second is a Gaussian process emulator, which treats $\log c_0(\eta)$ as an unknown latent function and uses a probabilistic interpolation model to predict its value at unobserved values of η . The third is a shape-constrained additive model (SCAM), which incorporates structural information about the normalising constant through monotonicity constraints.

Since

$$\frac{\partial}{\partial \eta} \log c_0(\eta) = \mathbb{E}_{\pi_\eta} [\log L(D_0 \mid \boldsymbol{\beta})], \quad (5.13)$$

the derivative process provides information about the qualitative shape of $\log c_0(\eta)$, which we incorporate through constrained smoothers. The convexity of $\eta \mapsto \log c_0(\eta)$,

established in Appendix A.4, provides a principled basis for shape-constrained emulators; monotonicity holds empirically in all experiments but is not proved in general.

These structural properties motivate the use of shape-constrained emulators in Section 3.1.3. In all cases, we emulate the log-normalising constant using $K = 40$ points sampled via importance sampling. Three discrepancy scenarios are generated by increasing the separation between the historical and current regression coefficients. As in the previous experiments, increasing discrepancy is expected to shift posterior mass toward smaller values of η , indicating reduced borrowing from historical information.

The Exchange sampler is treated as the reference method. For each surrogate model, we report Wasserstein distances relative to the Exchange posterior, together with effective sample sizes, acceptance rates, posterior means, posterior standard deviations, and computational runtime (including the time required for fitting the surrogate model). Posterior agreement is assessed for η and for each regression coefficient separately. This allows distortions introduced by the surrogate approximation to be localized to specific components of the parameter vector.




Figure 6 – Posterior density of η across discrepancy scenarios for the logistic regression normalised power prior experiment.




Figure 7 – Posterior density of the first regression coefficient across discrepancy scenarios.




Figure 8 – Posterior density of the second regression coefficient across discrepancy scenarios.




Figure 9 – Posterior density of the third regression coefficient across discrepancy scenarios.

Scenario	Method	W_η	W_{β_1}	W_{β_2}	W_{β_3}	ESS_η	Acc. η	Time (s)
Small	Exchange MH	0.000	0.000	0.000	0.000	56	0.947	3598.9
Small	Spline Emulator	0.074	0.017	0.018	0.015	33	0.955	1.1
Small	GP Emulator	0.070	0.019	0.017	0.013	33	0.956	7.5
Small	SCAM Emulator	0.106	0.018	0.022	0.027	39	0.955	19.6
Moderate	Exchange MH	0.000	0.000	0.000	0.000	33	0.918	3564.4
Moderate	Spline Emulator	0.087	0.052	0.016	0.029	26	0.921	1.1
Moderate	GP Emulator	0.145	0.080	0.019	0.016	29	0.920	6.5
Moderate	SCAM Emulator	0.014	0.017	0.011	0.015	37	0.938	15.8
Large	Exchange MH	0.000	0.000	0.000	0.000	309	0.604	3180.1
Large	Spline Emulator	0.002	0.024	0.030	0.021	411	0.700	1.1
Large	GP Emulator	0.002	0.061	0.028	0.020	397	0.682	5.5
Large	SCAM Emulator	0.013	0.123	0.016	0.066	72	0.716	14.2

Table 5 – Diagnostic comparison for the logistic regression NPP experiment. Wasserstein distances are computed against the exchange MH reference within each discrepancy scenario.

Scenario	Method	$\bar{\eta}$	$\text{sd}(\eta)$	$\text{MCSE}(\eta)$	β_1	$\text{sd}(\beta_1)$	$\text{MCSE}(\beta_1)$
Small	Exchange MH	0.681	0.227	0.0305	0.860	0.198	0.0081
Small	Spline Emulator	0.608	0.254	0.0442	0.875	0.189	0.0080
Small	GP Emulator	0.611	0.254	0.0443	0.874	0.187	0.0080
Small	SCAM Emulator	0.575	0.257	0.0410	0.878	0.200	0.0083
Moderate	Exchange MH	0.395	0.245	0.0426	1.437	0.270	0.0145
Moderate	Spline Emulator	0.308	0.229	0.0446	1.489	0.284	0.0146
Moderate	GP Emulator	0.250	0.177	0.0327	1.517	0.281	0.0140
Moderate	SCAM Emulator	0.396	0.234	0.0383	1.421	0.272	0.0190
Large	Exchange MH	0.059	0.044	0.0025	3.383	0.687	0.0418
Large	Spline Emulator	0.059	0.043	0.0021	3.370	0.675	0.0377
Large	GP Emulator	0.060	0.044	0.0022	3.444	0.753	0.0502
Large	SCAM Emulator	0.072	0.083	0.0098	3.259	0.645	0.0387

Table 6 – Posterior summaries for the logistic regression NPP experiment across three discrepancy levels.

5.2.3.1 Discussion of results

The results reveal an important asymmetry between the inferential impact of approximation errors on η and on β . Even in cases where the surrogate model introduces notable distortion in the marginal posterior of η , the posterior distributions of the regression coefficients β remain largely unaffected. This suggests that, for the quantity of primary inferential interest in regression applications, moderate errors in the approximation of $c_0(\eta)$ may be practically tolerable.

A further notable finding is that all methods produce the most discernible differences under the moderate-discrepancy scenario. Under large discrepancy, the posterior of η concentrates near zero regardless of the method, reducing the sensitivity to approximation errors. Under small discrepancy, all posteriors are diffuse and similar.

It is precisely the intermediate regime, where the posterior of η is neither degenerate nor flat, that discriminates most sharply between exact and approximate approaches. The spline and GP surrogates tend to concentrate posterior mass towards zero more aggressively than SCAM and Exchange in this regime, suggesting that they may underestimate the degree of historical borrowing when data are moderately compatible.

Finally, the Exchange sampler is considerably slower per iteration than the surrogate methods, reflecting the cost of exact auxiliary draws via Stan. This computational gap motivates the development of faster exact alternatives, such as the pseudo-marginal approach of Section 4.3, as a practical middle ground.

6 Discussion

6.1 Summary of contributions

This thesis has developed exact Markov chain Monte Carlo methods for posterior inference under the normalised power prior (NPP), addressing a computational gap that has persisted in the literature since the NPP was formalised by [Duan, Ye and Smith \(2006\)](#) and further studied by [Carvalho and Ibrahim \(2021\)](#). The central challenge is the doubly intractable structure of the NPP posterior: the prior itself contains a normalising constant $c_0(\eta)$ that depends on the power parameter η and cannot be evaluated in closed form for general likelihoods (Section 2.3). Previous approaches addressed this by approximating $c_0(\eta)$ through function emulation or Monte Carlo integration, but these methods lack formal guarantees and their bias is difficult to quantify without an exact reference.

Our main theoretical contribution is the derivation of the thermodynamic identity (2.31), which expresses the log-derivative of $c_0(\eta)$ as a tractable conditional expectation under the power prior. This identity provides the analytical foundation for two complementary exact algorithms. The Poisson estimator (Section 4.2.1) uses this identity to construct a non-negative unbiased estimator of the normalising constant ratio $c_0(\eta)/c_0(\eta')$, which can be embedded within a pseudo-marginal Metropolis–Hastings scheme while preserving exact invariance. The exchange algorithm (Section 4.1) avoids evaluating $c_0(\eta)$ altogether by introducing an auxiliary draw from the proposed power prior, cancelling the intractable ratio through a careful augmentation of the target space.

6.2 Empirical findings

The exchange algorithm is the most practically useful of the exact methods developed here. When exact simulation from the power prior $\pi_{\eta'}(\theta \mid D_0)$ is available, as it is for Gaussian and Normal–Inverse–Gamma likelihoods, the exchange update is both exact and computationally tractable, making it the natural choice for constructing reference posteriors against which other methods can be assessed. The experiments of Chapter 4 confirm that the exchange sampler faithfully recovers the analytic posterior across all discrepancy regimes, providing the first systematic exact baseline for the NPP in the literature.

The approximate methods studied in Chapter 3 (spline emulation, Gaussian process emulation, and shape-constrained additive models) all perform competitively with the exchange sampler when measured on the regression coefficients β , which are typically

the primary inferential target in applications. Distortions in the marginal posterior of η introduced by the surrogate approximations have limited downstream effect on β , suggesting that moderate approximation errors in $c_0(\eta)$ may be practically tolerable in many settings. Among the surrogate methods, the SCAM emulator stands out for its more conservative behaviour in the medium-discrepancy regime, where it better preserves posterior mass at intermediate values of η relative to unconstrained spline and GP emulators. This confirms that incorporating the structural properties of $\log c_0(\eta)$, established in Appendix A.4, into the surrogate model leads to more reliable inference.

6.3 Limitations and future directions

The exact methods developed here require either exact simulation from the power prior (exchange algorithm) or the availability of deterministic bounds on the historical log-likelihood (Poisson estimator). Both conditions impose non-trivial requirements on the model class. For the exchange algorithm, exact simulation from $\pi_{\eta'}(\theta \mid D_0)$ is available in conjugate families but must be replaced by an approximate inner MCMC chain in general non-conjugate models, at the cost of exactness. For the Poisson estimator, the bounded-support assumption that enables tight local envelopes is satisfied for compact priors and Gaussian likelihoods on bounded domains, but fails for unbounded parameter spaces without additional truncation. Extending both constructions to a broader class of models is the most immediate direction for future work.

A second direction concerns the relationship between the exact and approximate methods. The experiments demonstrate that surrogate-based methods can be computationally far more efficient than exact approaches when the cost of exact powered-prior simulation is high, as in the logistic regression experiment where auxiliary draws require running Stan at every iteration. Developing strategies that adaptively switch between exact and approximate methods, for instance, using a surrogate for the bulk of MCMC exploration and correcting with an occasional exact step, could provide a principled route to scalable inference with finite-sample guarantees.

Finally, the methods and theory developed here are currently available as research code. A priority for future work is the development of a public R package, to be submitted to CRAN, that provides a black-box interface to both approximate and exact MCMC algorithms for the normalised power prior. Such a package would make these methods accessible to practitioners in clinical trial design and other applied settings where the NPP has seen the most interest.

Code Availability The code for all experiments and methods described in this thesis is available at <https://github.com/adamesalles/npp-msc-thesis>.

References

- Adame, E.; Carvalho, L. M. Bayesian shape-constrained curve-fitting with gaussian processes: prior elicitation and computation. In: *Proceeding Series of the Brazilian Society of Computational and Applied Mathematics*. 2025. v. 11, n. 1, p. 1–2.
- Andrieu, C.; Roberts, G. O. The pseudo-marginal approach for efficient Monte Carlo computations. *The Annals of Statistics*, Institute of Mathematical Statistics, v. 37, n. 2, p. 697 – 725, 2009. Available on: <https://doi.org/10.1214/07-AOS574>.
- Barker, A. A. Monte Carlo calculations of the radial distribution functions for a proton-electron plasma. *Australian J. Phys.*, Vol: 18, 03 1965. Available on: <https://www.osti.gov/biblio/4621785>.
- Carpenter, B. et al. Stan: A probabilistic programming language. *Journal of Statistical Software*, Foundation for Open Access Statistic, v. 76, n. 1, 2017. ISSN 1548-7660. Available on: <http://dx.doi.org/10.18637/jss.v076.i01>.
- Carvalho, L. M.; Ibrahim, J. G. On the normalized power prior. *Statistics in Medicine*, v. 40, n. 24, p. 5251–5275, 2021. Available on: <https://onlinelibrary.wiley.com/doi/abs/10.1002/sim.9124>.
- Chen, M.-H.; Ibrahim, J. G. Power prior distributions for regression models. *Statistical Science*, Institute of Mathematical Statistics, v. 15, n. 1, fev. 2000. ISSN 0883-4237. Available on: <http://dx.doi.org/10.1214/ss/1009212673>.
- Duan, Y.; Ye, K.; Smith, E. P. Evaluating water quality using power priors to incorporate historical information. *Environmetrics*, Wiley, v. 17, n. 1, p. 95–106, 2006. ISSN 1099-095X. Available on: <http://dx.doi.org/10.1002/env.752>.
- Friel, N.; Pettitt, A. N. Marginal likelihood estimation via power posteriors. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, Oxford University Press (OUP), v. 70, n. 3, p. 589–607, abr. 2008. ISSN 1467-9868. Available on: <http://dx.doi.org/10.1111/j.1467-9868.2007.00650.x>.
- Gelfand, A. E.; Smith, A. F. M. Sampling-based approaches to calculating marginal densities. *Journal of the American Statistical Association*, Informa UK Limited, v. 85, n. 410, p. 398–409, jun. 1990. ISSN 1537-274X. Available on: <http://dx.doi.org/10.1080/01621459.1990.10476213>.
- Gelman, A. et al. *Bayesian Data Analysis, Third Edition*. Taylor & Francis, 2013. (Chapman & Hall/CRC Texts in Statistical Science). ISBN 9781439840955. Available on: <https://books.google.com.br/books?id=ZXL6AQAAQBAJ>.
- Geman, S.; Geman, D. Stochastic relaxation, gibbs distributions, and the bayesian restoration of images. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, Institute of Electrical and Electronics Engineers (IEEE), PAMI-6, n. 6, p. 721–741, nov. 1984. ISSN 0162-8828. Available on: <http://dx.doi.org/10.1109/TPAMI.1984.4767596>.

Gilks, W. R.; Best, N. G.; Tan, K. K. C. Adaptive rejection metropolis sampling within gibbs sampling. *Applied Statistics*, JSTOR, v. 44, n. 4, p. 455, 1995. ISSN 0035-9254. Available on: <http://dx.doi.org/10.2307/2986138>.

Gonçalves, F. B.; Łatuszyński, K.; Roberts, G. O. Barker's algorithm for Bayesian inference with intractable likelihoods. *Brazilian Journal of Probability and Statistics*, Brazilian Statistical Association, v. 31, n. 4, p. 732 – 745, 2017.

Han, Z.; Zhang, Q.; Wang, M.; Ye, K.; Chen, M.-H. On efficient posterior inference in normalized power prior Bayesian analysis. *Biometrical Journal*, v. 65, n. 5, p. 2200194, 2023. Available on: <https://onlinelibrary.wiley.com/doi/abs/10.1002/bimj.202200194>.

Hastie, T.; Tibshirani, R. Generalized additive models. *Statistical Science*, Institute of Mathematical Statistics, v. 1, n. 3, ago. 1986. ISSN 0883-4237. Available on: <http://dx.doi.org/10.1214/ss/1177013604>.

Hastings, W. K. Monte carlo sampling methods using markov chains and their applications. *Biometrika*, Oxford University Press (OUP), v. 57, n. 1, p. 97–109, abr. 1970. ISSN 0006-3444. Available on: <http://dx.doi.org/10.1093/biomet/57.1.97>.

Hobbs, B. P.; Carlin, B. P.; Mandrekar, S. J.; Sargent, D. J. Hierarchical commensurate and power prior models for adaptive incorporation of historical information in clinical trials. *Biometrics*, Oxford University Press (OUP), v. 67, n. 3, p. 1047–1056, mar. 2011. ISSN 0006-341X. Available on: <http://dx.doi.org/10.1111/j.1541-0420.2011.01564.x>.

Ibrahim, J. G.; Chen, M.-H. Power prior distributions for regression models. *Statistical Science*, Institute of Mathematical Statistics, v. 15, n. 1, p. 46–60, 2000. ISSN 08834237, 21688745. Available on: <http://www.jstor.org/stable/2676676>.

Ibrahim, J. G.; Chen, M.-H.; Gwon, Y.; Chen, F. The power prior: theory and applications. *Statistics in Medicine*, v. 34, n. 28, p. 3724–3749, 2015. Available on: <https://onlinelibrary.wiley.com/doi/abs/10.1002/sim.6728>.

Liang, F. A double metropolis–hastings sampler for spatial models with intractable normalizing constants. *Journal of Statistical Computation and Simulation*, Informa UK Limited, v. 80, n. 9, p. 1007–1022, out. 2009. ISSN 1563-5163. Available on: <http://dx.doi.org/10.1080/00949650902882162>.

Metropolis, N.; Rosenbluth, A. W.; Rosenbluth, M. N.; Teller, A. H.; Teller, E. Equation of state calculations by fast computing machines. *The Journal of Chemical Physics*, AIP Publishing, v. 21, n. 6, p. 1087–1092, jun. 1953. ISSN 1089-7690. Available on: <http://dx.doi.org/10.1063/1.1699114>.

Meyn, S.; Tweedie, R. L.; Glynn, P. W. *Markov Chains and Stochastic Stability*. Cambridge University Press, 2009. ISBN 9780511626630. Available on: <http://dx.doi.org/10.1017/CBO9780511626630>.

Murray, I.; Ghahramani, Z.; MacKay, D. *MCMC for doubly-intractable distributions*. 2012. Available on: <https://arxiv.org/abs/1206.6848>.

Neuenschwander, B.; Branson, M.; Spiegelhalter, D. J. A note on the power prior. *Statistics in Medicine*, v. 28, n. 28, p. 3562–3566, 2009. Available on: <https://onlinelibrary.wiley.com/doi/abs/10.1002/sim.3722>.

- Neuenschwander, B.; Capkun-Niggli, G.; Branson, M.; Spiegelhalter, D. J. Summarizing historical information on controls in clinical trials. *Clinical Trials*, SAGE Publications, v. 7, n. 1, p. 5–18, jan. 2010. ISSN 1740-7753. Available on: <http://dx.doi.org/10.1177/1740774509356002>.
- Park, J.; Haran, M. *Bayesian Inference in the Presence of Intractable Normalizing Functions*. arXiv, 2017. Available on: <https://arxiv.org/abs/1701.06619>.
- Park, J.; Haran, M. A function emulation approach for doubly intractable distributions. *Journal of Computational and Graphical Statistics*, Taylor & Francis, v. 29, n. 1, p. 66–77, 2020.
- Polasek, W. Handbook of markov chain monte carlo edited by steve brooks, andrew gelman, galin jones, xiao-li meng. *International Statistical Review*, Wiley, v. 80, n. 1, p. 184–185, abr. 2012. ISSN 1751-5823. Available on: http://dx.doi.org/10.1111/j.1751-5823.2011.00179_9.x.
- Pyra, N.; Wood, S. N. Shape constrained additive models. *Statistics and Computing*, Springer Science and Business Media LLC, v. 25, n. 3, p. 543–559, fev. 2015. ISSN 1573-1375. Available on: <http://dx.doi.org/10.1007/s11222-013-9448-7>.
- Rasmussen, C. E.; Williams, C. K. I. *Gaussian Processes for Machine Learning*. The MIT Press, 2006. ISBN 9780262256834. Available on: <http://dx.doi.org/10.7551/mitpress/3206.001.0001>.
- Riihimäki, J.; Vehtari, A. Gaussian processes with monotonicity information. *Journal of Machine Learning Research - Proceedings Track*, v. 9, p. 645–652, 01 2010.
- Robert, C. P.; Casella, G. *Monte Carlo Statistical Methods*. Springer New York, 2004. ISSN 1431-875X. ISBN 9781475741452. Available on: <http://dx.doi.org/10.1007/978-1-4757-4145-2>.
- Schmidli, H. et al. Robust meta-analytic-predictive priors in clinical trials with historical control information. *Biometrics*, Oxford University Press (OUP), v. 70, n. 4, p. 1023–1032, out. 2014. ISSN 1541-0420. Available on: <http://dx.doi.org/10.1111/biom.12242>.
- Spiegelhalter, D. J. Incorporating Bayesian Ideas into Health-Care Evaluation. *Statistical Science*, v. 19, n. 1, fev. 2004. ISSN 0883-4237.
- Stumpf-Fétizon, T.; Gonçalves, F. B. *Scalable Bernoulli factories for Bayesian inference with intractable likelihoods*. arXiv, 2025. ArXiv:2505.05438 [stat]. Available on: <http://arxiv.org/abs/2505.05438>.
- Vats, D.; Gonçalves, F. B.; Łatuszyński, K.; Roberts, G. O. Efficient Bernoulli factory Markov chain Monte Carlo for intractable posteriors. *Biometrika*, v. 109, n. 2, p. 369–385, 10 2021. ISSN 1464-3510. Available on: <https://doi.org/10.1093/biomet/asab031>.
- Wang, X.; Berger, J. O. Estimating shape constrained functions using gaussian processes. *SIAM/ASA Journal on Uncertainty Quantification*, v. 4, n. 1, p. 1–25, 2016.
- Wood, S. N. *Generalized Additive Models: An Introduction with R*. Chapman and Hall/CRC, 2017. ISBN 9781315370279. Available on: <http://dx.doi.org/10.1201/9781315370279>.

APPENDIX A – Proofs

A.1 Thermodynamic identity for the normalising constant

Proof. Under regularity conditions permitting differentiation under the integral sign, we have

$$\frac{dc_0(\eta)}{d\eta} = \frac{d}{d\eta} \int_{\Theta} L_0(D_0 | t)^\eta dP_0(t) \quad (\text{A.1})$$

$$= \int_{\Theta} \frac{d}{d\eta} L_0(D_0 | t)^\eta dP_0(t) \quad (\text{A.2})$$

$$= \int_{\Theta} \log L_0(D_0 | t) \cdot L_0(D_0 | t)^\eta dP_0(t) \quad (\text{A.3})$$

$$= c_0(\eta) \int_{\Theta} \log L_0(D_0 | t) \cdot \pi_\eta(t | D_0) dt \quad (\text{A.4})$$

$$= c_0(\eta) \mathbb{E}_{\pi_\eta}[\log L_0(D_0 | \theta)]. \quad (\text{A.5})$$

Since $c_0(\eta) > 0$ whenever the initial prior is proper, dividing both sides by $c_0(\eta)$ yields (2.31). $\square$

A.2 Unbiasedness of the Poisson estimator

Proof. We use the notation of Section 4.2.1: write $\lambda := \lambda_{\eta, \eta'}$, $c := c_{\eta, \eta'}$, and $\Delta := \Delta(\eta, \eta')$. Since $(U_j, \theta_j) \stackrel{\text{iid}}{\sim} \mathbf{P}_{\eta, \eta'}$ and $K \sim \text{Poisson}(\lambda)$ are drawn independently, the tower property gives

$$\mathbb{E}[\widehat{I}(\eta, \eta')] = \exp\{\lambda - c\} \mathbb{E}_K \left[\left(\mathbb{E}_{\mathbf{P}_{\eta, \eta'}} \left[\frac{c - Z_{\eta, \eta'}}{\lambda} \right] \right)^K \right] \quad (\text{A.6})$$

$$= \exp\{\lambda - c\} \mathbb{E}_K \left[\left(\frac{c - \Delta}{\lambda} \right)^K \right]. \quad (\text{A.7})$$

Applying the probability-generating function of a $\text{Poisson}(\lambda)$ random variable, $\mathbb{E}[r^K] = \exp\{\lambda(r - 1)\}$, with $r = (c - \Delta)/\lambda$, gives

$$\mathbb{E}[\widehat{I}(\eta, \eta')] = \exp\{\lambda - c\} \cdot \exp \left\{ \lambda \left(\frac{c - \Delta}{\lambda} - 1 \right) \right\} \quad (\text{A.8})$$

$$= \exp\{\lambda - c\} \cdot \exp\{c - \Delta - \lambda\} \quad (\text{A.9})$$

$$= \exp\{-\Delta\} = \frac{c_0(\eta)}{c_0(\eta')}, \quad (\text{A.10})$$

which establishes the claim. $\square$

A.3 Correctness of the exchange algorithm

Proof. We verify detailed balance for the augmented kernel. Let $\pi(\eta \mid \theta) \propto \pi_A(\eta) L_0(D_0 \mid \theta)^\eta / c_0(\eta)$ denote the conditional target for η . The exchange algorithm augments the state space with θ^* drawn from the proposed normalised model:

$$\theta^* \sim \pi_{\eta'}(\theta \mid D_0) = \frac{L_0(D_0 \mid \theta)^{\eta'}}{c_0(\eta')} \pi_0(\theta). \quad (\text{A.11})$$

The augmented target on (η, θ^*) is

$$\tilde{\pi}(\eta, \theta^*) \propto \pi(\eta \mid \theta) \cdot \pi_{\eta'}(\theta^* \mid D_0). \quad (\text{A.12})$$

The proposed move $(\eta, \theta^*) \rightarrow (\eta', \theta)$ is accepted with the Metropolis–Hastings ratio

$$R_{\text{ex}} = \frac{\tilde{\pi}(\eta', \theta) q(\eta \mid \eta')}{\tilde{\pi}(\eta, \theta^*) q(\eta' \mid \eta)} = \frac{\pi_A(\eta') L_0(D_0 \mid \theta)^{\eta'} / c_0(\eta')}{\pi_A(\eta) L_0(D_0 \mid \theta)^\eta / c_0(\eta)} \cdot \frac{L_0(D_0 \mid \theta)^\eta / c_0(\eta)}{L_0(D_0 \mid \theta^*)^{\eta'} / c_0(\eta')} \cdot \frac{q(\eta \mid \eta')}{q(\eta' \mid \eta)}. \quad (\text{A.13})$$

The factors $c_0(\eta)$ and $c_0(\eta')$ cancel, yielding

$$R_{\text{ex}} = \frac{\pi_A(\eta')}{\pi_A(\eta)} \cdot \frac{q(\eta \mid \eta')}{q(\eta' \mid \eta)} \cdot \frac{L_0(D_0 \mid \theta)^{\eta' - \eta}}{L_0(D_0 \mid \theta^*)^{\eta' - \eta}}, \quad (\text{A.14})$$

which is exactly the ratio in Section 4.1. The augmented kernel therefore satisfies detailed balance with respect to $\tilde{\pi}$. It remains to verify that the marginal of $\tilde{\pi}$ over θ^* equals the desired target $\pi(\eta \mid \theta)$. Indeed,

$$\int_{\Theta} \tilde{\pi}(\eta, \theta^*) d\theta^* \propto \pi(\eta \mid \theta) \int_{\Theta} \pi_{\eta'}(\theta^* \mid D_0) d\theta^* = \pi(\eta \mid \theta), \quad (\text{A.15})$$

since $\pi_{\eta'}(\cdot \mid D_0)$ is a proper probability density integrating to one. Hence the marginal invariant distribution of η under the augmented kernel is $\pi(\eta \mid \theta)$, as required. $\square$

A.4 Convexity of log-normalising constant

Proof. We show that $\eta \mapsto \log c_0(\eta)$ is convex on $[0, 1]$. By (2.31), the second derivative satisfies

$$\frac{d^2 \log c_0(\eta)}{d\eta^2} = \frac{d}{d\eta} \mathbb{E}_{\pi_\eta} [\log L_0(D_0 \mid \theta)] = \text{Var}_{\pi_\eta} [\log L_0(D_0 \mid \theta)] \geq 0, \quad (\text{A.16})$$

where the identity follows by differentiating under the integral sign and recognising the resulting expression as a variance under π_η . Since the second derivative is non-negative, $\log c_0(\eta)$ is convex. $\square$

APPENDIX B – Analytic Formulations

This appendix collects closed-form expressions for the normalising constant $c_0(\eta)$ and its derivative under specific model families. These results are used in the experiments of Chapter 4 to construct analytic reference chains and to validate the unbiased estimators developed there.

B.1 Gaussian likelihood with known variance

Consider $d_{0j} \sim \mathcal{N}(\theta, \sigma^2)$ with known variance σ^2 for all $j = 1, \dots, N_0$, and let P_0 be the measure induced by the prior $\pi_0(\theta) \propto 1$.

Therefore, let $\bar{d}_{0p} = \sum_{j=1}^p d_{0j}/p$ be the mean of the historical data D_0 and $S^2 = \sum_{j=1}^p (d_{0j} - \bar{d}_{0p})^2$ the sample variance. Then, we have

$$\begin{aligned}
 c_0(\eta) &= \int_{\Theta} L(D_0 | t)^\eta dP_0(t), \\
 &= \int_{-\infty}^{\infty} \frac{1}{(2\pi\sigma^2)^{\eta p/2}} \exp \left\{ -\frac{\eta}{2\sigma^2} \left(\sum_{j=1}^p (d_{0j} - t)^2 \right) \right\} dt, \\
 &= \frac{1}{(2\pi\sigma^2)^{\eta p/2}} \int_{-\infty}^{\infty} \exp \left\{ -\frac{\eta}{2\sigma^2} \left(\sum_{j=1}^p ((d_{0j} - \bar{d}_{0p}) + (\bar{d}_{0p} - t))^2 \right) \right\} dt, \\
 &= \frac{1}{(2\pi\sigma^2)^{\eta p/2}} \int_{-\infty}^{\infty} \exp \left\{ -\frac{\eta}{2\sigma^2} \left(S^2 + p(t - \bar{d}_{0p})^2 + \sum_{j=1}^p 2(d_{0j} - \bar{d}_{0p})(t - \bar{d}_{0p}) \right) \right\} dt, \\
 &= \frac{1}{(2\pi\sigma^2)^{\eta p/2}} \exp \left\{ -\frac{\eta}{2\sigma^2} S^2 \right\} \int_{-\infty}^{\infty} \exp \left\{ -\frac{\eta}{2\sigma^2} p(t - \bar{d}_{0p})^2 \right\} dt, \\
 &= \frac{1}{(2\pi\sigma^2)^{\eta p/2}} \exp \left\{ -\frac{\eta}{2\sigma^2} S^2 \right\} \sqrt{\frac{2\pi\sigma^2}{p\eta}} = \frac{1}{(2\pi\sigma^2)^{(\eta p-1)/2} \sqrt{p\eta}} \exp \left\{ -\frac{\eta}{2\sigma^2} S^2 \right\}.
 \end{aligned}$$

Computing the logarithm of the normalising constant, we have

$$\log c_0(\eta) = -\frac{(\eta p - 1)}{2} \log(2\pi\sigma^2) - \frac{1}{2} \log(p\eta) - \frac{\eta}{2\sigma^2} S^2.$$

Therefore, the derivative of the logarithm of the normalising constant is given by

$$\frac{d \log c_0(\eta)}{d\eta} = -\frac{p}{2} \log(2\pi\sigma^2) - \frac{S^2}{2\sigma^2} - \frac{1}{2\eta}.$$

On the other hand, we can show that $\mathbb{E}_{\pi_\eta}[\log(L(D_0 | \theta))]$ matches the previous expression. In this case, we have that $\theta | D_0 \sim \mathcal{N}(\bar{d}_{0p}, \sigma^2/p\eta)$, so we can compute the expectation as follows:

$$\begin{aligned}
\mathbb{E}_{\pi_\eta}[\log(L(D_0 | \theta))] &= -\frac{p}{2} \log(2\pi\sigma^2) - \mathbb{E}_{\pi_\eta} \left[\frac{1}{2\sigma^2} \left(\sum_{j=1}^p (d_{0j} - \theta)^2 \right) \right], \\
&= -\frac{p}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{j=1}^p \mathbb{E}_{\pi_\eta} [(d_{0j} - \theta)^2], \\
&= -\frac{p}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{j=1}^p ((d_{0j} - \mathbb{E}_{\pi_\eta}[\theta])^2 + \mathbb{V}_{\pi_\eta}[\theta]), \\
&= -\frac{p}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{j=1}^p \left((d_{0j} - \bar{d}_{0p})^2 + \frac{\sigma^2}{p\eta} \right), \\
&= -\frac{p}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{j=1}^p (d_{0j} - \bar{d}_{0p})^2 - \frac{p}{2\sigma^2} \frac{\sigma^2}{p\eta}, \\
&= -\frac{p}{2} \log(2\pi\sigma^2) - \frac{S^2}{2\sigma^2} - \frac{1}{2\eta}.
\end{aligned}$$

B.1.1 Adding a Gaussian initial prior

Consider again $d_{0j} \sim \mathcal{N}(\theta, \sigma^2)$ with known variance σ^2 for all $j = 1, \dots, N_0$, but now let $P_0 \equiv \mathcal{N}(0, \sigma_0^2)$, where σ_0^2 is a known initial prior variance.

Again, let $\bar{d}_{0p} = \sum_{j=1}^p d_{0j}/p$ be the mean of the historical data D_0 and $S^2 = \sum_{j=1}^p (d_{0j} - \bar{d}_{0p})^2$ the sample variance. Starting, we get:

$$c_0(\eta) = \frac{1}{(2\pi\sigma^2)^{\eta p/2}} \exp \left\{ -\frac{\eta}{2\sigma^2} S^2 \right\} \frac{1}{\sqrt{2\pi\sigma_0^2}} \int_{-\infty}^{\infty} \exp \left\{ -\frac{\eta}{2\sigma^2} p(t - \bar{d}_{0p})^2 - \frac{1}{2\sigma_0^2} t^2 \right\} dt.$$

Expanding $-\frac{\eta p}{2\sigma^2} (\bar{d}_{0p} - t)^2 - \frac{t^2}{2\sigma_0^2}$, we get $(\bar{d}_{0p} - t)^2 = t^2 - 2t\bar{d}_{0p} + \bar{d}_{0p}^2$. And, thus, we have that

$$\begin{aligned}
-\frac{\eta p}{2\sigma^2} (\bar{d}_{0p} - t)^2 - \frac{t^2}{2\sigma_0^2} &= -\frac{\eta p}{2\sigma^2} (t^2 - 2t\bar{d}_{0p} + \bar{d}_{0p}^2) - \frac{t^2}{2\sigma_0^2}, \\
&= -\frac{1}{2} \left(t^2 \left(\frac{\eta p}{\sigma^2} + \frac{1}{\sigma_0^2} \right) - 2t\bar{d}_{0p} \frac{\eta p}{\sigma^2} + \bar{d}_{0p}^2 \frac{\eta p}{\sigma^2} \right).
\end{aligned}$$

Define $A = \frac{\eta p}{\sigma^2} + \frac{1}{\sigma_0^2}$ and $B = \frac{\eta p}{\sigma^2} \bar{d}_{0p}$. We can now complete the square. Once we have that $At^2 - 2Bt = A \left(t - \frac{B}{A} \right)^2 - \frac{B^2}{A}$, it follows that $-\frac{1}{2}(At^2 - 2Bt) = -\frac{A}{2} \left(t - \frac{B}{A} \right)^2 + \frac{B^2}{2A}$. Which gives us

$$-\frac{1}{2} \left(At^2 - 2Bt + \bar{d}_{0p}^2 \frac{\eta p}{\sigma^2} \right) = -\frac{A}{2} \left(t - \frac{B}{A} \right)^2 + \frac{B^2}{2A} - \frac{\bar{d}_{0p}^2 \eta p}{2\sigma^2}.$$

Therefore, we can rewrite the normalising constant as

$$\begin{aligned}
c_0(\eta) &= \frac{1}{(2\pi\sigma^2)^{\eta p/2}} \exp\left\{-\frac{\eta}{2\sigma^2}S^2\right\} \frac{1}{\sqrt{2\pi\sigma_0^2}} \int_{-\infty}^{\infty} \exp\left\{-\frac{A}{2}\left(t - \frac{B}{A}\right)^2 + \frac{B^2}{2A} - \frac{\bar{d}_{0p}^2}{2}\frac{\eta p}{\sigma^2}\right\} dt, \\
&= \frac{1}{(2\pi\sigma^2)^{\eta p/2}} \exp\left\{-\frac{\eta}{2\sigma^2}S^2 - \frac{\eta p}{2\sigma^2}\bar{d}_{0p}^2\right\} \cdot \frac{1}{\sqrt{A\sigma_0^2}} \cdot \exp\left\{\frac{B^2}{2A}\right\}, \\
&= \frac{1}{(2\pi\sigma^2)^{\eta p/2}} \exp\left\{-\frac{\eta}{2\sigma^2}\left(S^2 - p\bar{d}_{0p}^2\right) + \frac{B^2}{2A}\right\} \cdot \frac{1}{\sqrt{A\sigma_0^2}}.
\end{aligned}$$

Taking the logarithm, we have

$$\log c_0(\eta) = -\frac{(\eta p)}{2} \log(2\pi\sigma^2) - \frac{\eta}{2\sigma^2} \left(S^2 - p\bar{d}_{0p}^2\right) + \frac{B^2}{2A} - \frac{1}{2} \log(A\sigma_0^2).$$

And, thus, we can compute the derivative of the logarithm of the normalising constant term by term. The first two terms and the last one can be calculated right away:

$$\begin{aligned}
\frac{d}{d\eta} \left(-\frac{(\eta p)}{2} \log(2\pi\sigma^2) \right) &= -\frac{p}{2} \log(2\pi\sigma^2), \\
\frac{d}{d\eta} \left(-\frac{\eta}{2\sigma^2} \left(S^2 - p\bar{d}_{0p}^2\right) \right) &= -\frac{1}{2\sigma^2} \left(S^2 - p\bar{d}_{0p}^2\right), \text{ and} \\
\frac{d}{d\eta} \left(-\frac{1}{2} \log(A\sigma_0^2) \right) &= -\frac{1}{2} \left(\frac{A'}{A} \right) = \frac{-p/\sigma^2}{2A}.
\end{aligned}$$

The third term is given by

$$\frac{d}{d\eta} \left(\frac{B^2}{2A} \right) = \frac{2BB'A - B^2A'}{2A^2},$$

where $B' = \frac{d}{d\eta} \left(\frac{\eta p}{\sigma^2} \bar{d}_{0p} \right) = \frac{p}{\sigma^2} \bar{d}_{0p}$ and $A' = \frac{d}{d\eta} \left(\frac{\eta p}{\sigma^2} + \frac{1}{\sigma_0^2} \right) = \frac{p}{\sigma^2}$, thus

$$\begin{aligned}
2BB' &= 2 \cdot \left(\frac{\eta p}{\sigma^2} \bar{d}_{0p} \right) \cdot \left(\frac{p}{\sigma^2} \bar{d}_{0p} \right) = 2\eta \frac{p^2}{\sigma^4} \bar{d}_{0p}^2, \\
B^2 &= \left(\frac{\eta p}{\sigma^2} \bar{d}_{0p} \right)^2 = \eta^2 \frac{p^2}{\sigma^4} \bar{d}_{0p}^2,
\end{aligned}$$

then

$$\frac{2BB'A - B^2A'}{2A^2} = \frac{1}{2A^2} \left(2\eta \frac{p^2}{\sigma^4} \bar{d}_{0p}^2 A - \eta^2 \frac{p^3}{\sigma^6} \bar{d}_{0p}^2 \right) = \frac{\eta \frac{p^2}{\sigma^4} \bar{d}_{0p}^2}{A^2} \left(A - \frac{\eta p}{2\sigma^2} \right).$$

And now we can recover the $c'_0(\eta)$ expression as

$$c'_0(\eta) = c_0(\eta)(\log(c_0(\eta)))' \quad (\text{B.1})$$

and compare it with the importance sampling reformulation.

B.2 Linear regression with NIG priors

Suppose we observe $\mathbf{y}_0 = \mathbf{X}_0\boldsymbol{\beta} + \varepsilon_0$, $\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \varepsilon$ with $\varepsilon_0 \sim \mathcal{N}(0, \sigma^2 I_{n_0})$, $\varepsilon \sim \mathcal{N}(0, \sigma^2 I_n)$ and unknown variance σ^2 . Now let $P_0 \equiv \mathcal{N}(\boldsymbol{\mu}_0, \mathbf{V}_0) \cdot \text{Inv-Gamma}(a, b)$, where $\boldsymbol{\mu}_0, \mathbf{V}_0, a$ and b are known hyperparameters.

Assume that $\mathbf{X}^\top \mathbf{X}$ is positive definite and let

$$\begin{aligned} \hat{\boldsymbol{\beta}}_0 &= (\mathbf{X}_0^\top \mathbf{X}_0)^{-1} \mathbf{X}_0^\top \mathbf{y}_0, \\ \hat{\boldsymbol{\beta}} &= (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}, \\ S_0 &= (\mathbf{y}_0 - \mathbf{X}_0 \hat{\boldsymbol{\beta}}_0)^\top (\mathbf{y}_0 - \mathbf{X}_0 \hat{\boldsymbol{\beta}}_0), \\ S &= (\mathbf{y} - \mathbf{X} \hat{\boldsymbol{\beta}})^\top (\mathbf{y} - \mathbf{X} \hat{\boldsymbol{\beta}}). \end{aligned}$$

So the likelihood of the historical data D_0 is

$$L(D_0 \mid \boldsymbol{\beta}, \sigma^2) = (2\pi\sigma^2)^{-n_0/2} \exp \left\{ -\frac{1}{2\sigma^2} [S_0 + (\boldsymbol{\beta} - \hat{\boldsymbol{\beta}}_0)^\top \mathbf{X}_0^\top \mathbf{X}_0 (\boldsymbol{\beta} - \hat{\boldsymbol{\beta}}_0)] \right\}$$

and the initial prior of $(\boldsymbol{\beta}, \sigma^2)$ is

$$\pi_0(\boldsymbol{\beta}, \sigma^2) \propto (\sigma^2)^{-a-p/2} \exp \left\{ -\frac{1}{\sigma^2} \left[b + \frac{1}{2} (\boldsymbol{\beta} - \boldsymbol{\mu}_0)^\top \mathbf{V}_0^{-1} (\boldsymbol{\beta} - \boldsymbol{\mu}_0) \right] \right\}.$$

Then

$$\begin{aligned} c_0(\eta) &= \int_0^\infty \int_{\mathbb{R}^p} \pi_0(\mathbf{u}, v) L(D_0 \mid \mathbf{u}, v)^\eta d\mathbf{u} dv, \\ &\propto \int_0^\infty \int_{\mathbb{R}^p} (2\pi)^{-n_0\eta/2} v^{-a-p/2-n_0\eta/2} \\ &\quad \cdot \exp \left\{ -\frac{1}{2v} [2H_0(\eta) + (\mathbf{u} - \tilde{\boldsymbol{\beta}}_0)^\top (\mathbf{V}_0^{-1} + \eta \mathbf{X}_0^\top \mathbf{X}_0) (\mathbf{u} - \tilde{\boldsymbol{\beta}}_0)] \right\} d\mathbf{u} dv, \\ &\propto (2\pi)^{-n_0\eta/2} \Gamma(\nu_0) \mid \Lambda_0 \mid^{-1/2} H_0(\eta)^{-\nu_0}, \end{aligned}$$

where

$$\begin{aligned}\tilde{\boldsymbol{\beta}}_0 &= (\Lambda_0)^{-1}(\mathbf{V}_0^{-1}\boldsymbol{\mu}_0 + \mathbf{X}_0^\top \mathbf{y}_0 \eta), \\ \nu_0 &= \frac{n_0 \eta}{2} + a - 1, \\ \Lambda_0 &= \mathbf{V}_0^{-1} + \mathbf{X}_0^\top \mathbf{X}_0 \eta, \\ H_0(\eta) &= b + \frac{\eta}{2} \left[S_0 + (\boldsymbol{\mu}_0 - \hat{\boldsymbol{\beta}}_0)^\top \mathbf{X}_0^\top \mathbf{X}_0 (\Lambda_0)^{-1} \mathbf{V}_0^{-1} (\boldsymbol{\mu}_0 - \hat{\boldsymbol{\beta}}_0) \right].\end{aligned}$$

This leads to the following NPP:

$$\pi(\boldsymbol{\beta}, \sigma^2, \eta \mid D_0) \propto \frac{\pi_0(\eta) H_0(\eta)^{\nu_0} |\Lambda_0|^{1/2}}{(\sigma^2)^{\frac{n+n_0\eta+p}{2}+a} \Gamma(\nu_0)} \exp \left\{ -\frac{1}{2\sigma^2} [(\boldsymbol{\beta} - \tilde{\boldsymbol{\beta}}_0)^\top \Lambda_0 (\boldsymbol{\beta} - \tilde{\boldsymbol{\beta}}_0) + 2H_0(\eta)] \right\}.$$

Including the likelihood we can derive the joint posterior distribution:

$$\pi(\boldsymbol{\beta}, \sigma^2, \eta \mid D, D_0) \propto \frac{\pi_0(\eta) H_0(\eta)^{\nu_0} |\Lambda_0|^{1/2}}{(\sigma^2)^{\frac{n+n_0\eta+p}{2}+a} \Gamma(\nu_0)} \exp \left\{ -\frac{1}{2\sigma^2} [(\boldsymbol{\beta} - \tilde{\boldsymbol{\beta}})^\top \Lambda (\boldsymbol{\beta} - \tilde{\boldsymbol{\beta}}) + 2H(\eta)] \right\},$$

where

$$\begin{aligned}\tilde{\boldsymbol{\beta}} &= \Lambda^{-1}(\mathbf{V}_0^{-1}\boldsymbol{\mu}_0 + \eta \mathbf{X}_0^\top \mathbf{y}_0 + \mathbf{X}^\top \mathbf{y}), \\ \Lambda &= \Lambda_0 + \mathbf{X}^\top \mathbf{X}, \\ H(\eta) &= H_0(\eta) + \frac{1}{2} [S + (\tilde{\boldsymbol{\beta}}_0 - \hat{\boldsymbol{\beta}})^\top \mathbf{X}^\top \mathbf{X} \Lambda^{-1} \Lambda_0 (\tilde{\boldsymbol{\beta}}_0 - \hat{\boldsymbol{\beta}})].\end{aligned}$$

So we can express the marginal posterior distribution of η given (D_0, D) as

$$\pi(\eta \mid D_0, D) \propto \frac{\pi_0(\eta) \Gamma(n/2 + \nu_0) |\Lambda_0|^{1/2} H_0(\eta)^{\nu_0}}{|\Lambda|^{1/2} \Gamma(\nu_0) H(\eta)^{\nu_0+n/2}},$$

and the conditional posterior distribution of $\boldsymbol{\beta}$ given (η, D_0, D) follows the multivariate Student distribution with location parameter $\tilde{\boldsymbol{\beta}}$, shape matrix Σ , and degree of freedom ρ , where $\rho = n + 2\nu_0$ and $\Sigma = \frac{2H(\eta)}{n+2\nu_0} \Lambda^{-1}$.